

CS162 Operating Systems and Systems Programming Lecture 4

Processes (con't), Threads, Concurrency

January 30th, 2020
Prof. John Kubiatowicz
<http://cs162.eecs.Berkeley.edu>

Recall: Modern Process with Threads

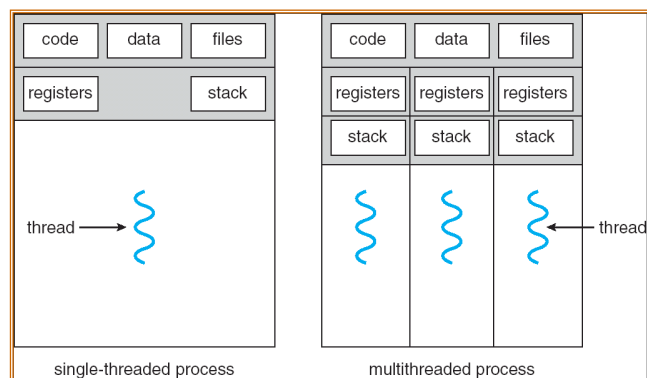
- Process: **execution environment with restricted rights**
 - Address Space with One or More Threads
 - » *One Page table per process!*
 - Owns memory (mapped pages)
 - Owns file descriptors, file system context, ...
 - Encapsulates one or more threads sharing process resources
- Thread: *a sequential execution stream within process*
(Sometimes called a “**Lightweight process**”)
 - Process still contains a single Address Space
 - No protection between threads
- Multithreading: *a single program made up of a number of different concurrent activities*
 - Sometimes called multitasking, as in Ada ...
- Why separate the concept of a thread from that of a process?
 - Discuss the “thread” part of a process (concurrency)
 - Separate from the “address space” (protection)
 - Heavyweight Process \equiv Process with one thread

1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.2

Recall: Single and Multithreaded Processes



- Threads encapsulate concurrency: “Active” component
- Address spaces encapsulate protection: “Passive” part
 - Keeps buggy program from trashing the system
- Why have multiple threads per address space?


1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.3

Recall: How do we Multiplex Processes?

- The current state of process held in a process control block (PCB):
 - This is a “snapshot” of the execution and protection environment
 - Only one PCB active at a time
- Give out CPU time to different processes (Scheduling):
 - Only one process “running” at a time
 - Give more time to important processes
- Give pieces of resources to different processes (Protection):
 - Controlled access to non-CPU resources
 - Example mechanisms:
 - » Memory Translation: Give each process their own (protected) address space
 - » Kernel/User duality: Arbitrary multiplexing of I/O through system calls



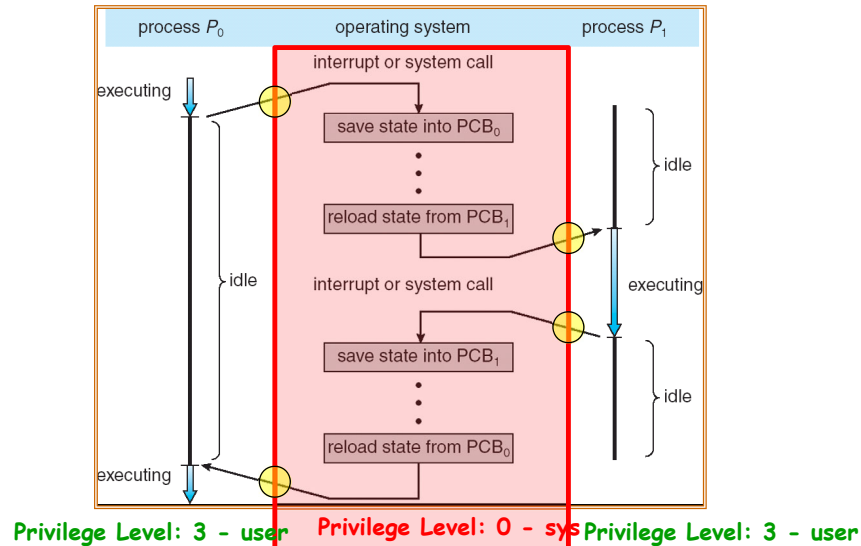
Process
Control
Block

1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.4

Recall: Context Switch

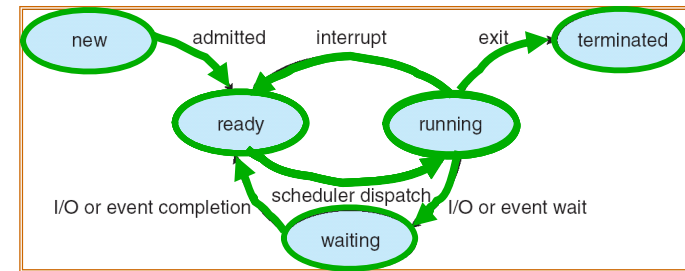


1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.5

Recall: Lifecycle of a Process



- As a process executes, it changes state:
 - new**: The process is being created
 - ready**: The process is waiting to run
 - running**: Instructions are being executed
 - waiting**: Process waiting for some event to occur
 - terminated**: The process has finished execution

1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.6

Discussion

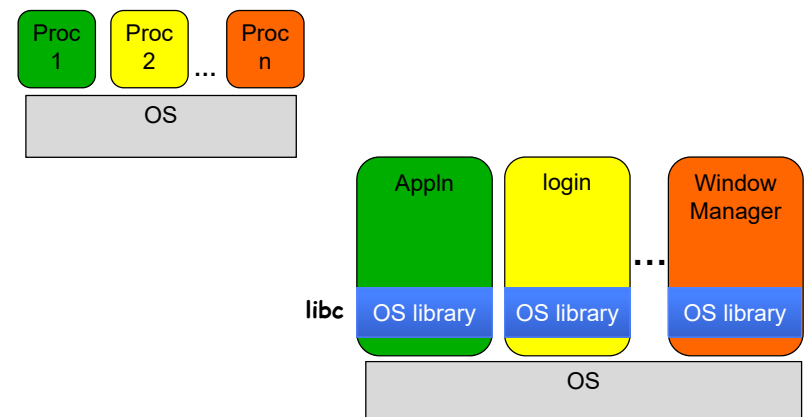
- Process is an *instance* of an *executing* program
 - The fundamental OS responsibility
 - Each instance has an identity (Process ID) or PID
- Processes do their work by processing and calling file system operations
 - This involves interacting with the Kernel!
 - How do we do that?
- Are there any operations on processes themselves?
 - create (fork) ?
 - terminate (exit) ?
 - sleep (sleep) ?
 - communicate with (e.g. signal)?

1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.7

OS Run-Time Library

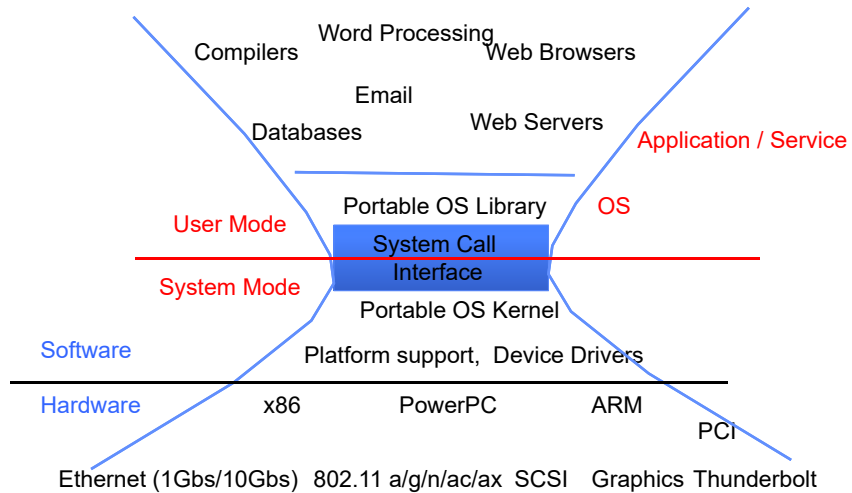


1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.8

A Narrow Waist



1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.9

pid.c

```
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <unistd.h>
#include <sys/types.h>
int main(int argc, char *argv[])
{
    pid_t pid = getpid(); /* get current processes PID */

    printf("My pid: %d\n", pid);

    exit(0);
}
```

ps anyone?

1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.10

POSIX/Unix

- **Portable Operating System Interface [X?]**
- Defines “Unix”, derived from AT&T Unix
 - Created to bring order to many Unix-derived OSs
- Interface for **application programmers** (mostly)

1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.11


System Calls

Application:

```
fd = open(pathname);
```

Library:

```
File *open(pathname) {
    asm code ... syscall # into ax
    put args into registers bx, ...
    special trap instruction
```



Operating System:
 get args from regs
 dispatch to system func
 process, schedule, ...
 complete, resume process

```
get results from regs
```

```
};
```

Continue with results

Pintos: userprog/syscall.c, lib/user/syscall.c

1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.12

SYSCALLs (of over 300)

%eax	Name	Source	%ebx	%ecx	%edx	%esi	%edi
1	sys_exit	kernel/exit.c	int	-	-	-	-
2	sys_fork	arch/i386/kernel/process.c	struct pt_regs	-	-	-	-
3	sys_read	fs/read_write.c	unsigned int	char *	size_t	-	-
4	sys_write	fs/read_write.c	unsigned int	const char *	size_t	-	-
5	sys_open	fs/open.c	const char *	int	int	-	-
6	sys_close	fs/open.c	unsigned int	-	-	-	-
7	sys_waitpid	kernel/exit.c	pid_t	unsigned int *	int	-	-
8	sys_creat	fs/open.c	const char *	int	-	-	-
9	sys_link	fs/namei.c	const char *	const char *	-	-	-
10	sys_unlink	fs/namei.c	const char *	-	-	-	-
11	sys_execve	arch/i386/kernel/process.c	struct pt_regs	-	-	-	-
12	sys_chdir	fs/open.c	const char *	-	-	-	-
13	sys_time	kernel/time.c	int *	-	-	-	-
14	sys_mknod	fs/namei.c	const char *	int	dev_t	-	-
15	sys_chmod	fs/open.c	const char *	mode_t	-	-	-
16	sys_lchown	fs/open.c	const char *	uid_t	gid_t	-	-
18	sys_stat	fs/stat.c	char *	struct _old_kernel_stat *	-	-	-
19	sys_lseek	fs/read_write.c	unsigned int	off_t	unsigned int	-	-
20	sys_getpid	kernel/sched.c	-	-	-	-	-
21	sys_mount	fs/super.c	char *	char *	char *	-	-
22	sys_oldumount	fs/super.c	char *	-	-	-	-
23	sys_setuid	kernel/sys.c	uid_t	-	-	-	-
24	sys_getuid	kernel/sched.c	-	-	-	-	-
25	sys_stime	kernel/time.c	int *	-	-	-	-
26	sys_ptrace	arch/i386/kernel/ptrace.c	long	long	long	long	-
27	sys_alarm	kernel/sched.c	unsigned int	-	-	-	-
28	sys_fstat	fs/stat.c	unsigned int	struct _old_kernel_stat *	-	-	-
29	sys_pause	arch/i386/kernel/sys_i386.c	-	-	-	-	-
30	sys_ftime	fs/open.c	char *	struct utimbuf *	-	-	-

Pintos: syscall-nr.h

Recall: Kernel System Call Handler

- Locate arguments
 - In registers or on user(!) stack
- Copy arguments
 - From user memory into kernel memory
 - Protect kernel from malicious code evading checks
- Validate arguments
 - Protect kernel from errors in user code
- Copy results back
 - into user memory

Process Management

- exit – terminate a process
- fork – copy the current process
- exec – change the *program* being run by the current process
- wait – wait for a process to finish
- kill – send a *signal* (interrupt-like notification) to another process
- sigaction – set handlers for signals

Process Management

- exit – terminate a process
- **fork – copy the current process**
- exec – change the *program* being run by the current process
- wait – wait for a process to finish
- kill – send a *signal* (interrupt-like notification) to another process
- sigaction – set handlers for signals

Creating Processes

- `pid_t fork();` -- copy the current process
 - This means everything!
 - New process has different pid
- Return value from **fork()**: pid (like an integer)
 - When > 0:
 - » Running in (original) **Parent** process
 - » return value is **pid** of new child
 - When = 0:
 - » Running in new **Child** process
 - When < 0:
 - » Error! Must handle somehow
 - » Running in original process
- If no error: State of original process duplicated in **both** Parent and Child!
 - Address Space (Memory), File Descriptors (covered later), etc...
 - Not as bad as it seems – really only copy page table [more later]

1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.17

fork1.c

```
#include <stdlib.h>
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>

int main(int argc, char *argv[]) {
    pid_t cpid, mypid;
    pid_t pid = getpid();          /* get current processes PID */
    printf("Parent pid: %d\n", pid);
    cpid = fork();
    if (cpid > 0) {                 /* Parent Process */
        mypid = getpid();
        printf("[%d] parent of [%d]\n", mypid, cpid);
    } else if (cpid == 0) {        /* Child Process */
        mypid = getpid();
        printf("[%d] child\n", mypid);
    } else {
        perror("Fork failed");
    }
}
```

1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.18

fork1.c

```
#include <stdlib.h>
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>

int main(int argc, char *argv[]) {
    pid_t cpid, mypid;
    pid_t pid = getpid();          /* get current processes PID */
    printf("Parent pid: %d\n", pid);
    cpid = fork();
    if (cpid > 0) {                 /* Parent Process */
        mypid = getpid();
        printf("[%d] parent of [%d]\n", mypid, cpid);
    } else if (cpid == 0) {        /* Child Process */
        mypid = getpid();
        printf("[%d] child\n", mypid);
    } else {
        perror("Fork failed");
    }
}
```

1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.19

fork1.c

```
#include <stdlib.h>
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>

int main(int argc, char *argv[]) {
    pid_t cpid, mypid;
    pid_t pid = getpid();          /* get current processes PID */
    printf("Parent pid: %d\n", pid);
    cpid = fork();
    if (cpid > 0) {                 /* Parent Process */
        mypid = getpid();
        printf("[%d] parent of [%d]\n", mypid, cpid);
    } else if (cpid == 0) {        /* Child Process */
        mypid = getpid();
        printf("[%d] child\n", mypid);
    } else {
        perror("Fork failed");
    }
}
```

1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.20

fork_race.c

```
int i;
cpid = fork();
if (cpid > 0) {
    for (i = 0; i < 10; i++) {
        printf("Parent: %d\n", i);
        // sleep(1);
    }
} else if (cpid == 0) {
    for (i = 0; i > -10; i--) {
        printf("Child: %d\n", i);
        // sleep(1);
    }
}
```

- What does this print?
- Would adding the calls to sleep matter?

Fork “race”

```
int i;
cpid = fork();
if (cpid > 0) {
    for (i = 0; i < 10; i++) {
        printf("Parent: %d\n", i);
        // sleep(1);
    }
} else if (cpid == 0) {
    for (i = 0; i > -10; i--) {
        printf("Child: %d\n", i);
        // sleep(1);
    }
}
```



Process Management

- fork – copy the current process
- exec – change the *program* being run by the current process
- **wait – wait for a process to finish**
- kill – send a *signal* (interrupt-like notification) to another process
- sigaction – set handlers for signals

fork2.c – parent waits for child to finish

```
int status;
pid_t tcpid;
...
cpid = fork();
if (cpid > 0) { /* Parent Process */
    mypid = getpid();
    printf("[%d] parent of [%d]\n", mypid, cpid);
    tcpid = wait(&status);
    printf("[%d] bye %d(%d)\n", mypid, tcpid, status);
} else if (cpid == 0) { /* Child Process */
    mypid = getpid();
    printf("[%d] child\n", mypid);
    exit(42);
}
...
```

Process Management

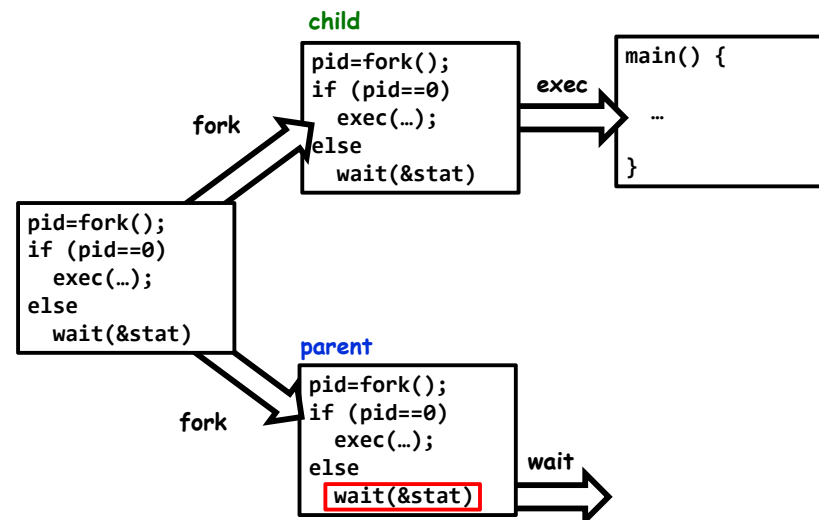
- fork – copy the current process
- **exec – change the *program* being run by the current process**
- wait – wait for a process to finish
- kill – send a *signal* (interrupt-like notification) to another process
- sigaction – set handlers for signals

1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.25

Process Management



1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.26

fork3.c

```
...
cpid = fork();
if (cpid > 0) {                /* Parent Process */
    tcpid = wait(&status);
} else if (cpid == 0) {        /* Child Process */
    char *args[] = {"ls", "-l", NULL};
    execv("/bin/ls", args);
    /* execv doesn't return when it works.
       So, if we got here, it failed! */
    perror("execv");
    exit(1);
}
...
```

1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.27

Shell

- A shell is a job control system
 - Allows programmer to create and manage a set of programs to do some task
 - Windows, MacOS, Linux all have shells
- Example: to compile a C program

```
cc -c sourcefile1.c
cc -c sourcefile2.c
ln -o program sourcefile1.o sourcefile2.o
./program
```



1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.28

Process Management

- fork – copy the current process
- exec – change the *program* being run by the current process
- wait – wait for a process to finish
- kill – send a *signal* (interrupt-like notification) to another process
- **sigaction – set handlers for signals**

inf_loop.c

```
#include <stdlib.h>
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
#include <signal.h>

void signal_callback_handler(int signum) {
    printf("Caught signal!\n");
    exit(1);
}

int main() {
    struct sigaction sa;
    sa.sa_flags = 0;
    sigemptyset(&sa.sa_mask);
    sa.sa_handler = signal_callback_handler;

    sigaction(SIGINT, &sa, NULL);
    while (1) {}
}
```

Common POSIX Signals

- SIGINT – control-C
- SIGTERM – default for kill shell command
- SIGSTP – control-Z (default action: stop process)
- SIGKILL, SIGSTOP – terminate/stop process
 - Can't be changed or disabled with sigaction
 - Why?

Administrivia

- Kubiatowicz Office Hours:
 - 1-2pm, Monday & Thursday
- No one left on WaitList or Concurrent Enrollment List!
 - Everyone who wants to be is in the class.
- **TOMORROW (Friday) is Drop Deadline!**
VERY HARD TO DROP LATER!
- Recommendation: Read assigned readings *before* lecture
- You should be going to sections now – Important information covered in section
 - Any section will do until groups assigned
- Group sign up via autograder form next week
 - Get finding groups of 4 people ASAP
 - Priority for same section; if cannot make this work, keep same TA
 - Remember: Your TA needs to see you in section!
- Midterm 1 conflicts
 - We will handle these conflicts next week

Reminder: Definitions

- A *thread* is a single execution sequence that represents a separately schedulable task
- Protection is an orthogonal concept
 - Can have one or many threads per protection domain
 - Single threaded user program: one thread, one protection domain
 - Multi-threaded user program: multiple threads, sharing same data structures, isolated from other user programs
 - Multi-threaded kernel: multiple threads, sharing kernel data structures, capable of using privileged instructions

1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.33

Threads Motivation

- Operating systems need to be able to handle *multiple things at once* (MTAO)
 - processes, interrupts, background system maintenance
- Servers need to handle MTAO
 - Multiple connections handled simultaneously
- Parallel programs need to handle MTAO
 - To achieve better performance
- Programs with user interfaces often need to handle MTAO
 - To achieve user responsiveness while doing computation
- Network and disk bound programs need to handle MTAO
 - To hide network/disk latency
 - Sequence steps in access or communication

1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.34

Silly Example for Threads

Imagine the following program:

```
main() {  
    ComputePI("pi.txt");  
    PrintClassList("classlist.txt");  
}
```

- What is the behavior here?
 - Program would never print out class list
 - Why? ComputePI would never finish

1/30/20


Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.35

Adding Threads

- Version of program with Threads (loose syntax):

```
main() {  
    thread_fork(ComputePI, "pi.txt" );  
    thread_fork(PrintClassList, "classlist.txt");  
}
```
- `thread_fork`: Start independent thread running given procedure
- What is the behavior here?
 - Now, you would actually see the class list
 - This *should* behave as if there are two separate CPUs



1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.36

More Practical Motivation

Back to Jeff Dean's "Numbers everyone should know":

Handle I/O in
separate
thread, avoid
blocking other
progress

L1 cache reference	0.5 ns
Branch mispredict	5 ns
L2 cache reference	7 ns
Mutex lock/unlock	25 ns
Main memory reference	100 ns
Compress 1K bytes with Zippy	3,000 ns
Send 2K bytes over 1 Gbps network	20,000 ns
Read 1 MB sequentially from memory	250,000 ns
Round trip within same datacenter	500,000 ns
Disk seek	10,000,000 ns
Read 1 MB sequentially from disk	20,000,000 ns
Send packet CA->Netherlands->CA	150,000,000 ns

Little Better Example for Threads?

Imagine the following program:

```
main() {  
    ...  
    ReadLargeFile("pi.txt");  
    RenderUserInterface();  
}
```

- What is the behavior here?
 - Still respond to user input
 - While reading file in the background


Voluntarily Giving Up Control

- I/O – e.g. keypress
- Waiting for a signal from another thread
 - Thread makes system call to *wait*
- Thread executes `thread_yield()`
 - Relinquishes CPU but puts calling thread back on ready queue

Adding Threads

- Version of program with Threads (loose syntax):

```
main() {  
    thread_fork(ReadLargeFile, "pi.txt" );  
    thread_fork(RenderUserInterface, "classlist.txt");  
}
```
- `thread_fork`: Start independent thread running given procedure
- What is the behavior here?
 - Now, you would actually see the class list
 - This *should* behave as if there are two separate CPUs



Thread State

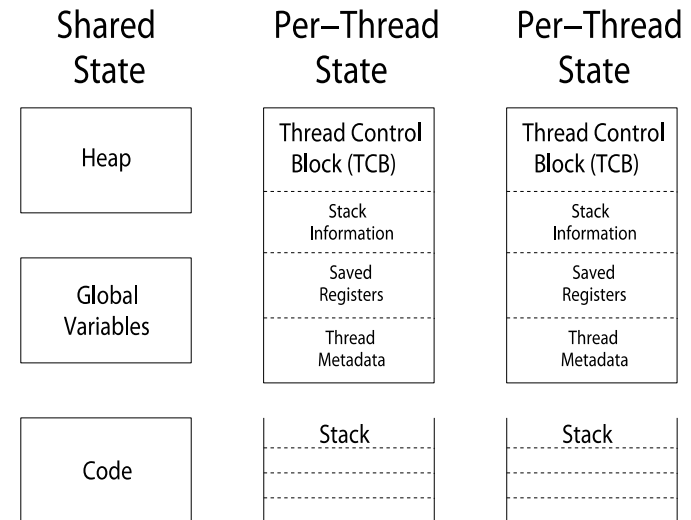
- State shared by all threads in process/address space
 - Content of memory (global variables, heap)
 - I/O state (file descriptors, network connections, etc)
- State “private” to each thread
 - Kept in **TCB** \equiv **Thread Control Block**
 - CPU registers (including, program counter)
 - Execution stack – what is this?
- Execution Stack
 - Parameters, temporary variables
 - Return PCs are kept while called procedures are executing

1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.41

Shared vs. Per-Thread State



1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.42

Execution Stack Example

```

A(int tmp) {
A:   if (tmp<2)
A+1:   B();
A+2:   printf(tmp);
      }
      B() {
B:     C();
B+1:   }
      C() {
C:     A(2);
C+1:   }
      A(1);
exit:
    
```

- Stack holds temporary results
- Permits recursive execution
- Crucial to modern languages

1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.43

Execution Stack Example

```

A(int tmp) {
A:   if (tmp<2)
A+1:   B();
A+2:   printf(tmp);
      }
      B() {
B:     C();
B+1:   }
      C() {
C:     A(2);
C+1:   }
      A(1);
exit:
    
```

Stack
Pointer

A: tmp=1
ret=exit

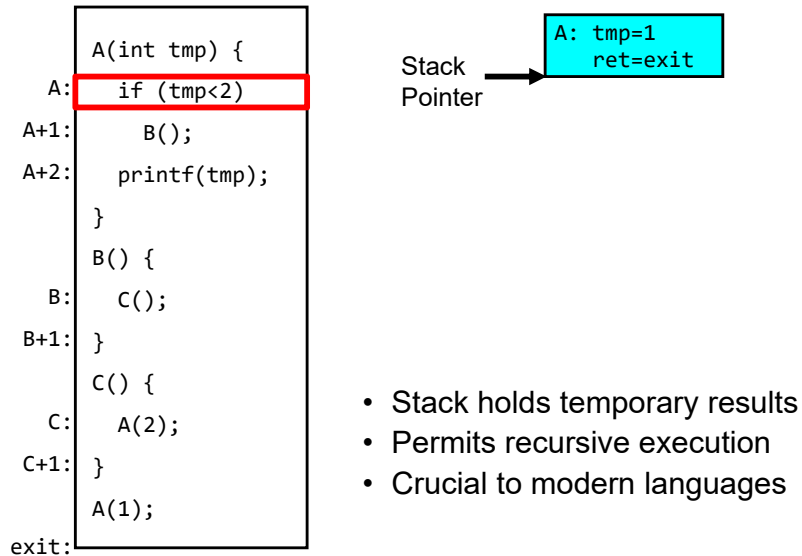
- Stack holds temporary results
- Permits recursive execution
- Crucial to modern languages

1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.44

Execution Stack Example

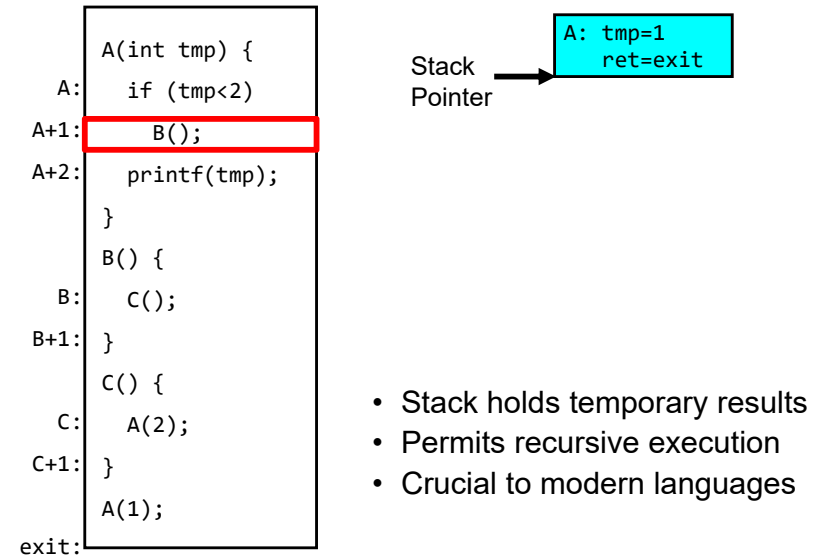


1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.45

Execution Stack Example

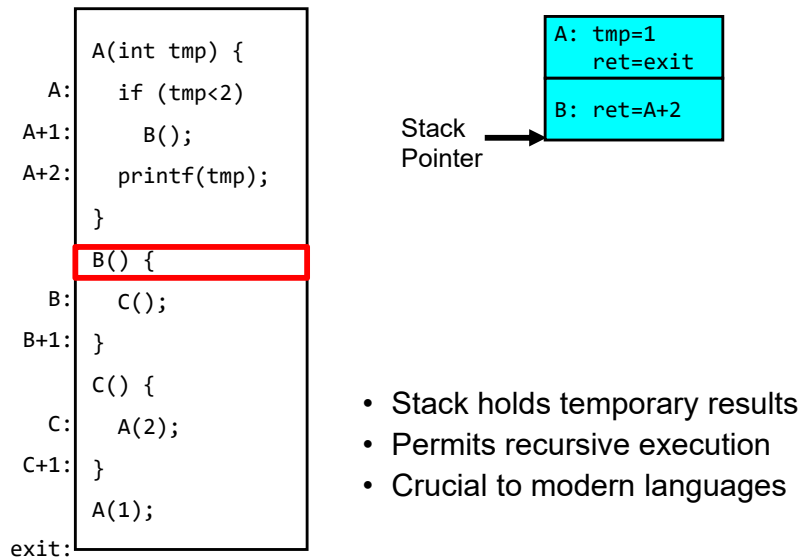


1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.46

Execution Stack Example

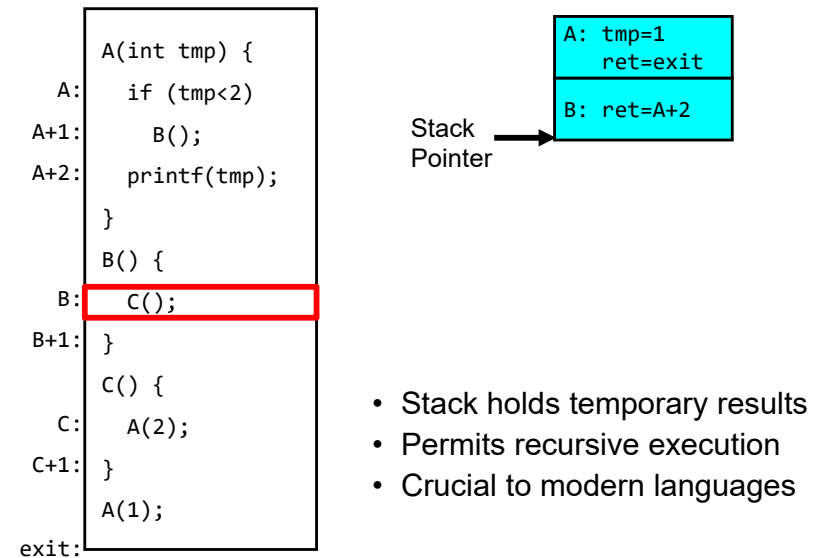


1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.47

Execution Stack Example

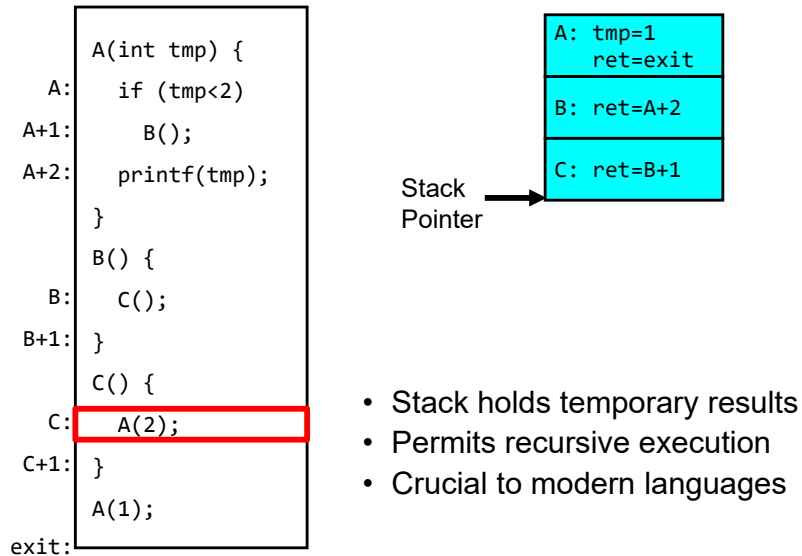


1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.48

Execution Stack Example

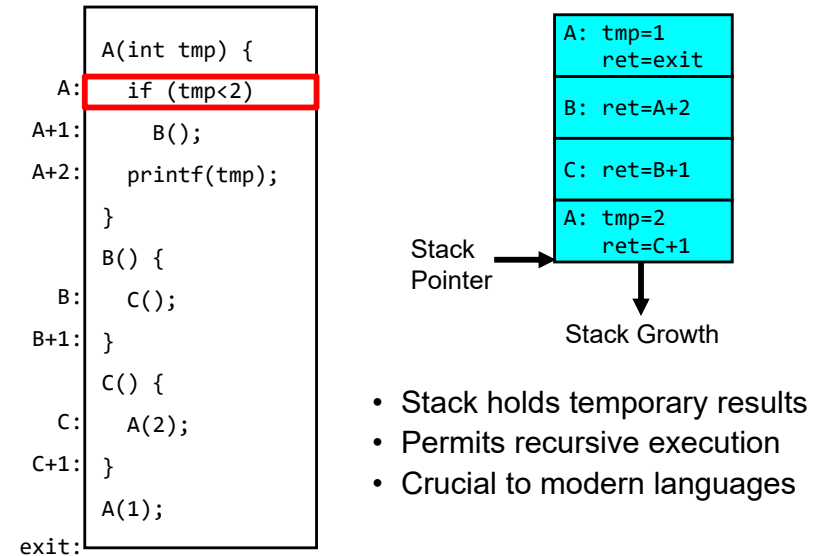


1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.49

Execution Stack Example

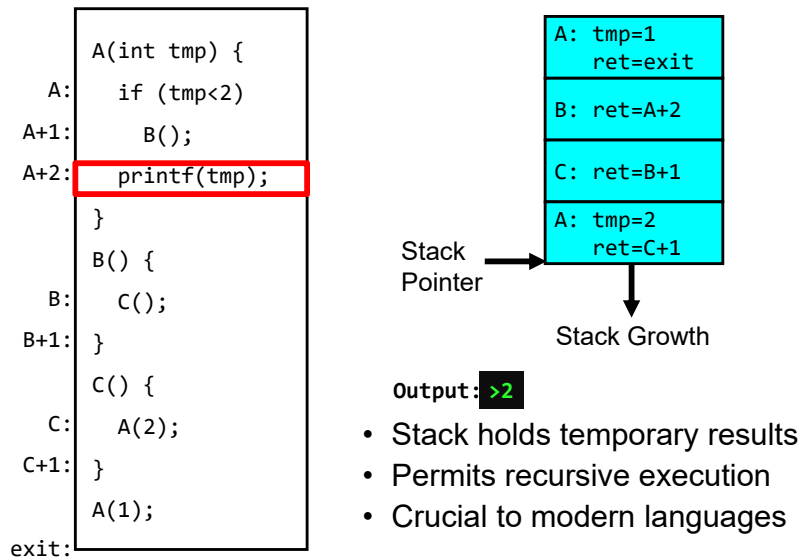


1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.50

Execution Stack Example

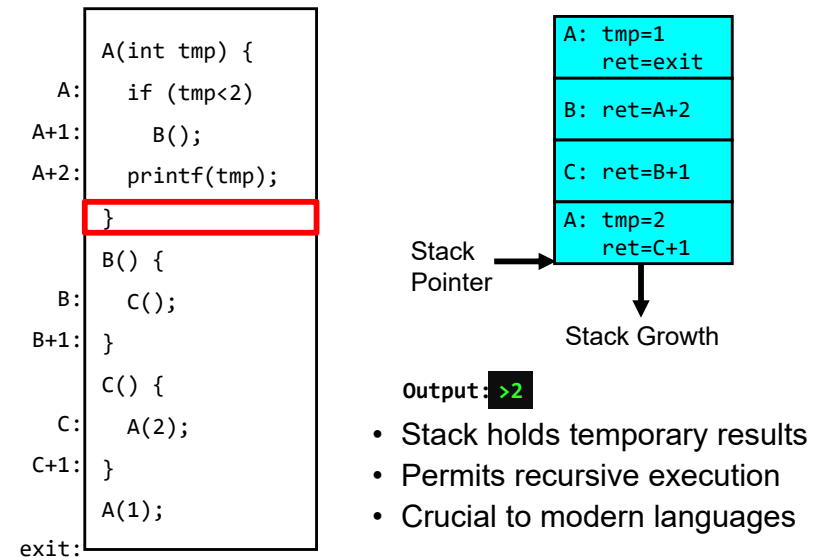


1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.51

Execution Stack Example

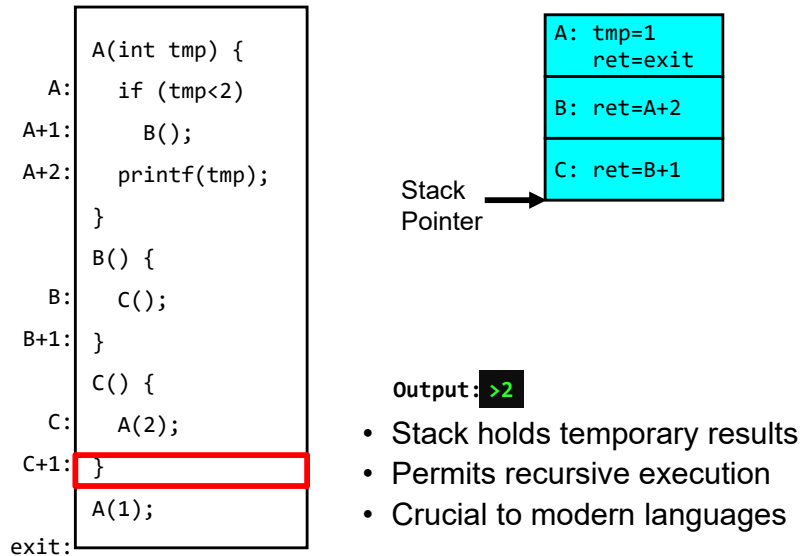


1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.52

Execution Stack Example

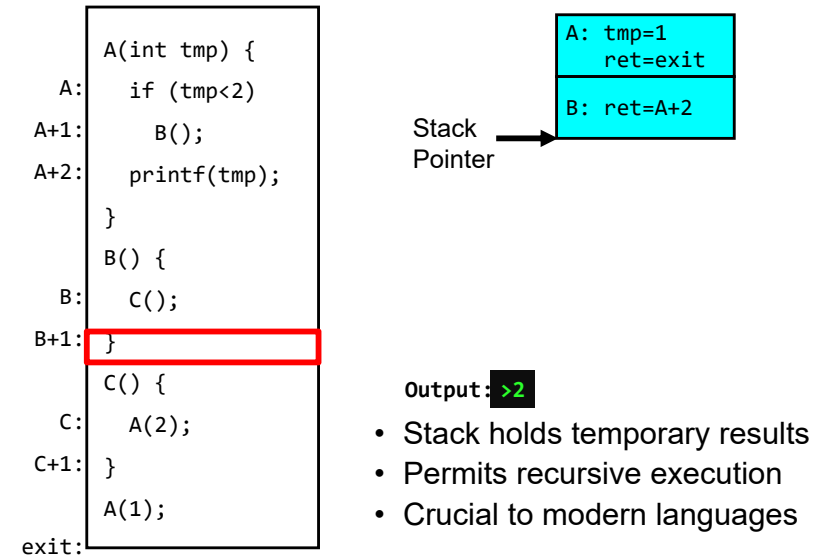


1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.53

Execution Stack Example

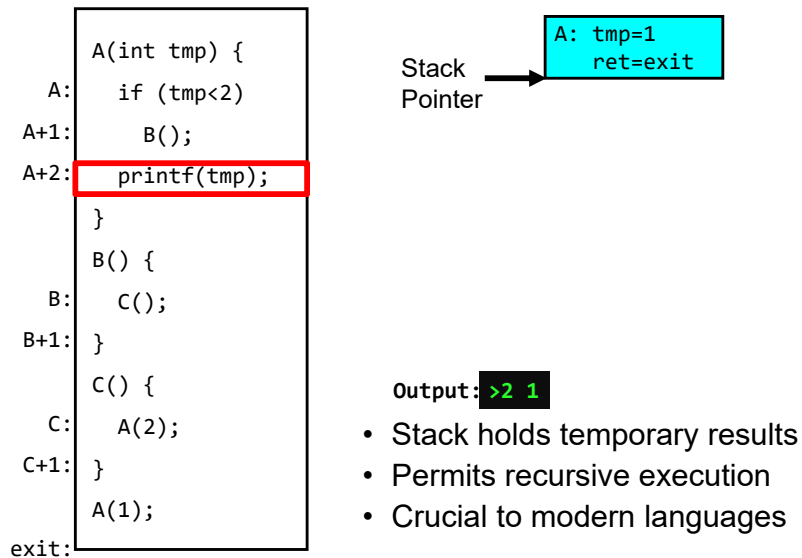


1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.54

Execution Stack Example

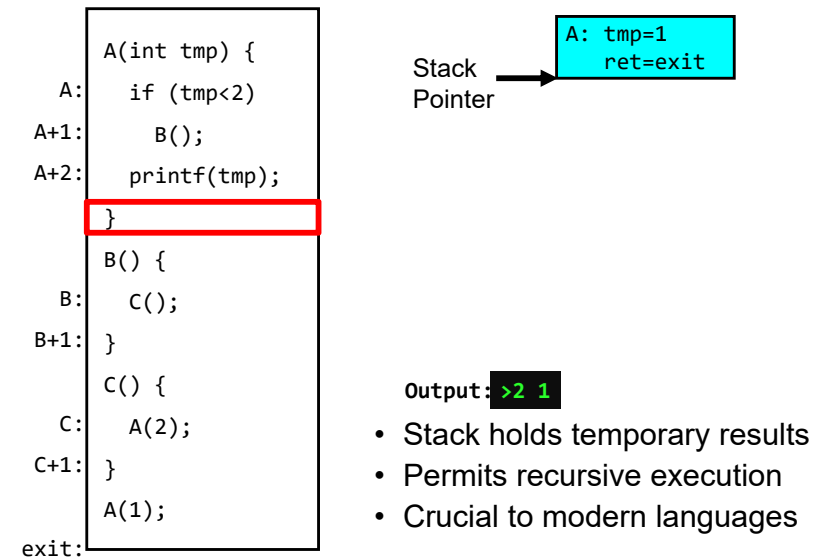


1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.55

Execution Stack Example



1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.56

Execution Stack Example

```
A(int tmp) {
    if (tmp<2)
        B();
    printf(tmp);
}
B() {
    C();
}
C() {
    A(2);
}
A(1);
```


Output: **>2 1**

- Stack holds temporary results
- Permits recursive execution
- Crucial to modern languages

Execution Stack Example

```
A(int tmp) {
    if (tmp<2)
        B();
    printf(tmp);
}
B() {
    C();
}
C() {
    A(2);
}
A(1);
```

Stack
Pointer



Stack Growth

- Stack holds temporary results
- Permits recursive execution
- Crucial to modern languages


Memory Footprint: Two-Threads

- If we stopped this program and examined it with a debugger, we would see

- Two sets of CPU registers
- Two sets of Stacks

- Questions:

- How do we position stacks relative to each other?
- What maximum size should we choose for the stacks?
- What happens if threads violate this?
- How might you catch violations?



Actual Thread Operations

- `thread_fork(func, args)`
 - Create a new thread to run `func(args)`
 - Pintos: `thread_create`
- `thread_yield()`
 - Relinquish processor voluntarily
 - Pintos: `thread_yield`
- `thread_join(thread)`
 - In parent, wait for forked thread to exit, then return
 - Pintos: `thread_join`
- `thread_exit`
 - Quit thread and clean up, wake up joiner if any
 - Pintos: `thread_exit`
- pThreads: POSIX standard for thread programming
[POSIX.1c, Threads extensions (IEEE Std 1003.1c-1995)]

Dispatch Loop

- Conceptually, the dispatching loop of the operating system looks as follows:

```
Loop {
    RunThread();
    ChooseNextThread();
    SaveStateOfCPU(curTCB);
    LoadStateOfCPU(newTCB);
}
```

- This is an *infinite* loop
 - One could argue that this is all that the OS does
- Should we ever exit this loop???
 - When would that be?

Running a thread

Consider first portion: RunThread()


- How do I run a thread?
 - Load its state (registers, PC, stack pointer) into CPU
 - Load environment (virtual memory space, etc)
 - Jump to the PC
- How does the dispatcher get control back?
 - Internal events: thread returns control voluntarily
 - External events: thread gets *preempted*

Internal Events

- Blocking on I/O
 - The act of requesting I/O implicitly yields the CPU
- Waiting on a “signal” from other thread
 - Thread asks to wait and thus yields the CPU
- Thread executes a yield()
 - Thread volunteers to give up CPU

```
computePI() {
    while(TRUE) {
        ComputeNextDigit();
        yield();
    }
}
```

Stack for Yielding Thread



- How do we run a new thread?

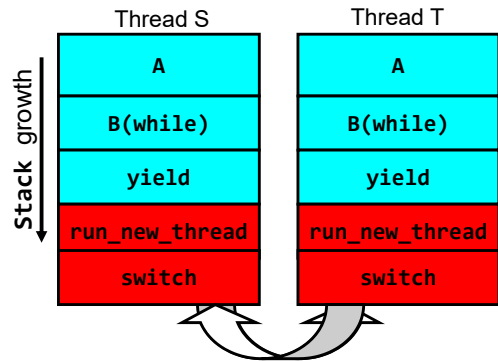
```
run_new_thread() {
    newThread = PickNewThread();
    switch(curThread, newThread);
    ThreadHouseKeeping(); /* Do any cleanup */
}
```
- How does dispatcher switch to a new thread?
 - Save anything next thread may trash: PC, regs, stack pointer
 - Maintain isolation for each thread

What Do the Stacks Look Like?

- Consider the following code blocks:

```
proc A() {  
    B();  
}  
proc B() {  
    while(TRUE) {  
        yield();  
    }  
}
```

- Suppose we have 2 threads:
 - Threads S and T



Thread S's switch returns to Thread T's (and vice versa)

Saving/Restoring state (often called "Context Switch")

```
Switch(tCur, tNew) {  
    /* Unload old thread */  
    TCB[tCur].regs.r7 = CPU.r7;  
    ...  
    TCB[tCur].regs.r0 = CPU.r0;  
    TCB[tCur].regs.sp = CPU.sp;  
    TCB[tCur].regs.retpc = CPU.retpc; /*return addr*/  
  
    /* Load and execute new thread */  
    CPU.r7 = TCB[tNew].regs.r7;  
    ...  
    CPU.r0 = TCB[tNew].regs.r0;  
    CPU.sp = TCB[tNew].regs.sp;  
    CPU.retpc = TCB[tNew].regs.retpc;  
    return; /* Return to CPU.retpc */  
}
```

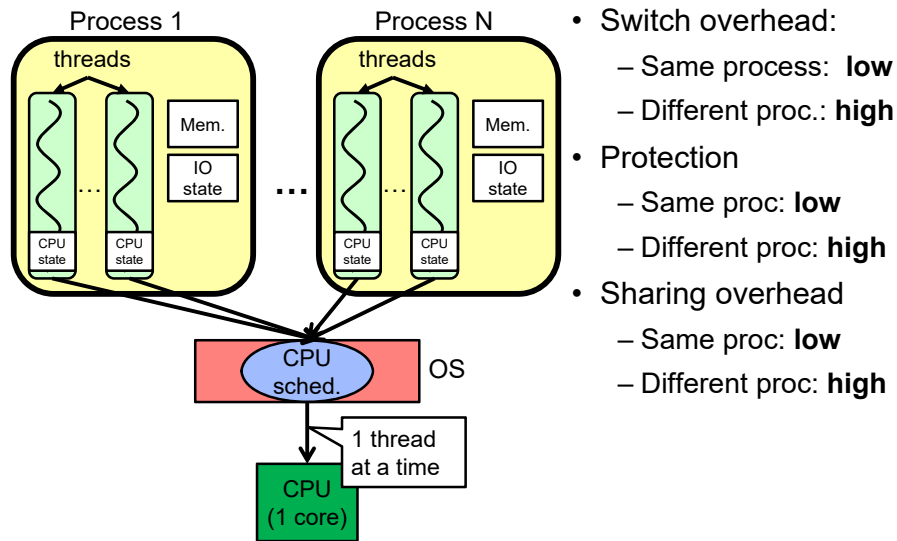
Switch Details (continued)

- What if you make a mistake in implementing switch?
 - Suppose you forget to save/restore register 32
 - Get intermittent failures depending on when context switch occurred and whether new thread uses register 32
 - System will give wrong result without warning
- Can you devise an exhaustive test to test switch code?
 - No! Too many combinations and inter-leavings
- Cautionary tale:
 - For speed, Topaz kernel saved one instruction in switch()
 - Carefully documented! Only works as long as kernel size < 1MB
 - What happened?
 - » Time passed, People forgot
 - » Later, they added features to kernel (no one removes features!)
 - » Very weird behavior started happening
 - Moral of story: Design for simplicity

Aren't we still switching contexts?

- Yes, but **much cheaper** than switching processes
 - No need to change address space
- Some numbers from Linux:
 - Frequency of context switch: 10-100ms
 - Switching between processes: 3-4 µsec.
 - Switching between threads: 100 ns

Processes vs. Threads

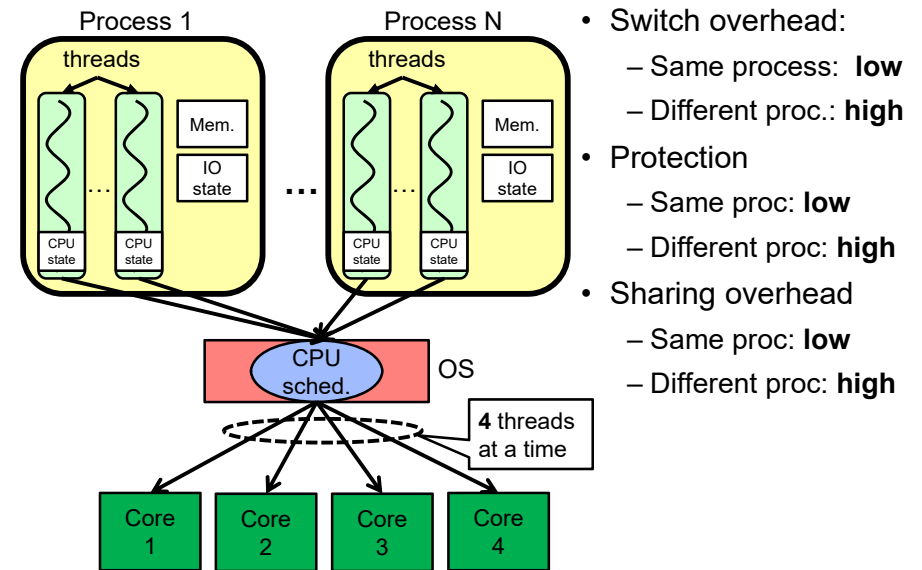


1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.69

Processes vs. Threads




1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.70

What happens when thread blocks on I/O?



- What happens when a thread requests a block of data from the file system?
 - User code invokes a system call
 - Read operation is initiated
 - Run new thread/switch
- Thread communication similar
 - Wait for Signal/Join
 - Networking

1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.71

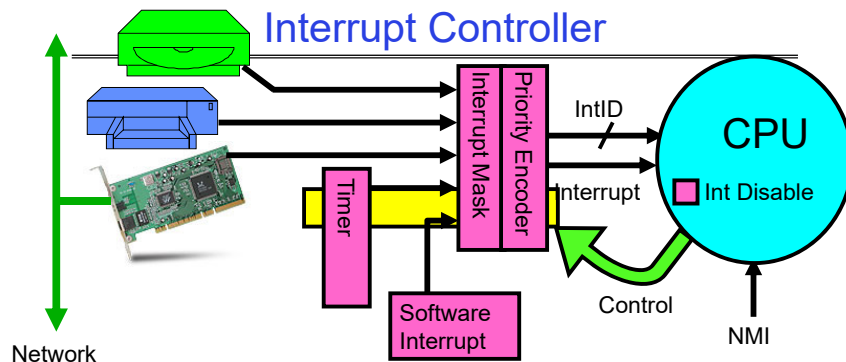
External Events

- What happens if thread never does any I/O, never waits, and never yields control?
 - Could the ComputePI program grab all resources and never release the processor?
 - » What if it didn't print to console?
 - Must find way that dispatcher can regain control!
- Answer: utilize external events
 - Interrupts: signals from hardware or software that stop the running code and jump to kernel
 - Timer: like an alarm clock that goes off every some milliseconds
- If we make sure that external events occur frequently enough, can ensure dispatcher runs

1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.72



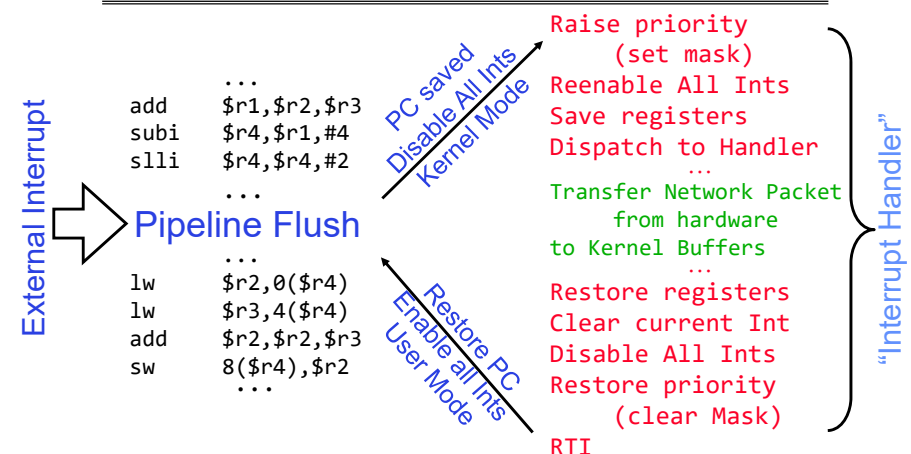
- Interrupts invoked with interrupt lines from devices
- Interrupt controller chooses interrupt request to honor
 - Interrupt identity specified with ID line
 - Mask enables/disables interrupts
 - Priority encoder picks highest enabled interrupt
 - Software Interrupt Set/Cleared by Software
- CPU can disable all interrupts with internal flag
- Non-Maskable Interrupt line (NMI) can't be disabled

1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.73

Example: Network Interrupt



- An interrupt is a hardware-invoked context switch
 - No separate step to choose what to run next
 - Always run the interrupt handler immediately


1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.74

Use of Timer Interrupt to Return Control

- Solution to our dispatcher problem
 - Use the timer interrupt to force scheduling decisions



- Timer Interrupt routine:

```
TimerInterrupt() {
    DoPeriodicHouseKeeping();
    run_new_thread();
}
```

1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.75

Hardware context switch support in x86

- Syscall/Intr (U → K)
 - PL 3 → 0;
 - TSS ← EFLAGS, CS:EIP;
 - SS:SP ← k-thread stack (TSS PL 0);
 - push (old) SS:ESP onto (new) k-stack
 - push (old) eflags, cs:eip, <err>
 - CS:EIP ← <k target handler>
- Then
 - Handler then saves other regs, etc
 - Does all its work, possibly choosing other threads, changing PTBR (CR3)
 - kernel thread has set up user GPRs
- iret (K → U)
 - PL 0 → 3;
 - Eflags, CS:EIP ← popped off k-stack
 - SS:SP ← user thread stack (TSS PL 3);




Figure 7-1. Structure of a Task

pg 2,942 of 4,922 of x86 reference manual

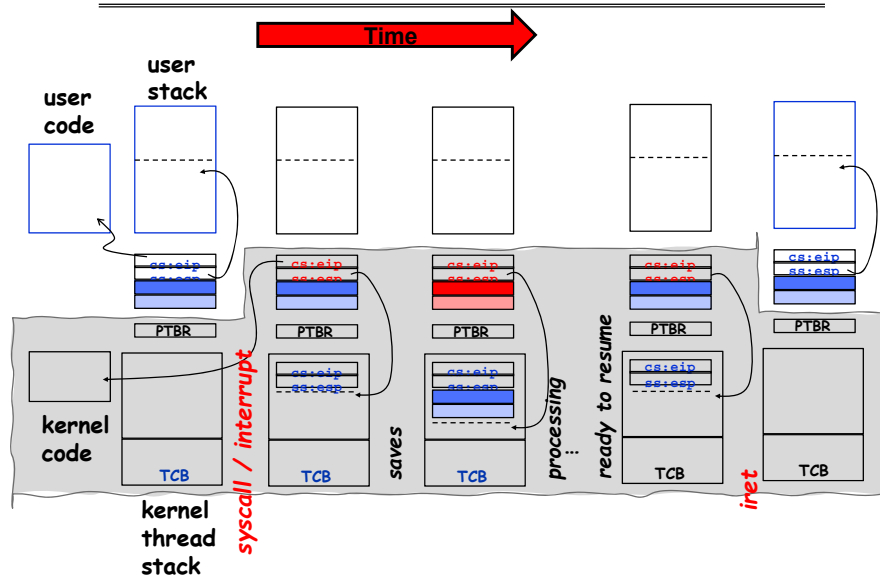
Pintos: tss.c, intr-stubs.S

1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.76

Pintos: Kernel Crossing on Syscall or Interrupt

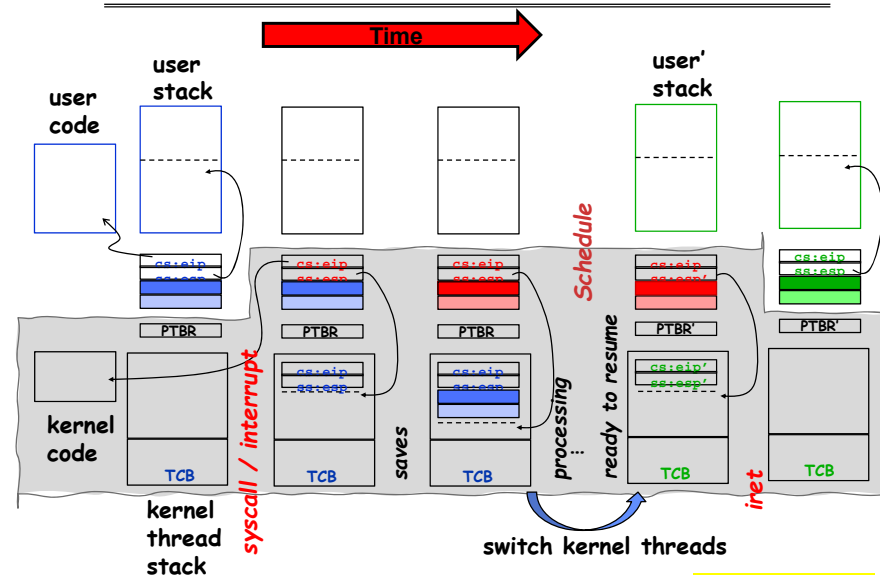


1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.77

Pintos: Context Switch – Scheduling



1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Pintos: switch.S

Lec 4.78

ThreadFork(): Create a New Thread

- ThreadFork() is a user-level procedure that creates a new thread and places it on ready queue
- Arguments to ThreadFork()
 - Pointer to application routine (fcnPtr)
 - Pointer to array of arguments (fcnArgPtr)
 - Size of stack to allocate
- Implementation
 - Sanity check arguments
 - Enter Kernel-mode and Sanity Check arguments again
 - Allocate new Stack and TCB
 - Initialize TCB and place on ready list (Runnable)

1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.79

How do we initialize TCB and Stack?

- Initialize Register fields of TCB
 - Stack pointer made to point at stack
 - PC return address \Rightarrow OS (asm) routine ThreadRoot()
 - Two arg registers (a0 and a1) initialized to fcnPtr and fcnArgPtr, respectively
- Initialize stack data?
 - No. Important part of stack frame is in registers (ra)
 - Think of stack frame as just before body of ThreadRoot() really gets started

ThreadRoot stub

Stack growth


Initial Stack

1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.80

How does Thread get started?




- Eventually, `run_new_thread()` will select this TCB and return into beginning of `ThreadRoot()`
 - This really starts the new thread

What does ThreadRoot() look like?

- `ThreadRoot()` is the root for the thread routine:


```
ThreadRoot() {
    DoStartupHousekeeping();
    UserModeSwitch(); /* enter user mode */
    Call fcnPtr(fcnArgPtr);
    ThreadFinish();
}
```

- Startup Housekeeping
 - Includes things like recording start time of thread
 - Other statistics
- Stack will grow and shrink with execution of thread
- Final return from thread returns into `ThreadRoot()` which calls `ThreadFinish()`
 - `ThreadFinish()` wake up sleeping threads



Multiprocessing vs Multiprogramming

- Remember Definitions:
 - Multiprocessing \equiv Multiple CPUs
 - Multiprogramming \equiv Multiple Jobs or Processes
 - Multithreading \equiv Multiple threads per Process
- What does it mean to run two threads "concurrently"?
 - Scheduler is free to run threads in any order and interleaving: FIFO, Random, ...
 - Dispatcher can choose to run each thread to completion or time-slice in big chunks or small chunks



Correctness for systems with concurrent threads

- If dispatcher can schedule threads in any way, programs must work under all circumstances
 - Can you test for this?
 - How can you know if your program works?
- Independent Threads:**
 - No state shared with other threads
 - Deterministic \Rightarrow Input state determines results
 - Reproducible \Rightarrow Can recreate Starting Conditions, I/O
 - Scheduling order doesn't matter (if `switch()` works!!!)
- Cooperating Threads:**
 - Shared State between multiple threads
 - Non-deterministic
 - Non-reproducible
- Non-deterministic and Non-reproducible means that bugs can be intermittent
 - Sometimes called "Heisenbugs"

Interactions Complicate Debugging

- Is any program truly independent?
 - Every process shares the file system, OS resources, network, etc
 - Extreme example: buggy device driver causes thread A to crash “independent thread” B
- You probably don’t realize how much you depend on reproducibility:
 - Example: Evil C compiler
 - » Modifies files behind your back by inserting errors into C program unless you insert debugging code
 - Example: Debugging statements can overrun stack
- Non-deterministic errors are really difficult to find
 - Example: Memory layout of kernel+user programs
 - » depends on scheduling, which depends on timer/other things
 - » Original UNIX had a bunch of non-deterministic errors
 - Example: Something which does interesting I/O
 - » User typing of letters used to help generate secure keys

1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.85

Why allow cooperating threads?


- People cooperate; computers help/enhance people’s lives, so computers must cooperate
 - By analogy, the non-reproducibility/non-determinism of people is a notable problem for “carefully laid plans”
- Advantage 1: Share resources
 - One computer, many users
 - One bank balance, many ATMs
 - » What if ATMs were only updated at night?
 - Embedded systems (robot control: coordinate arm & hand)
- Advantage 2: Speedup
 - Overlap I/O and computation
 - » Many different file systems do read-ahead
 - Multiprocessors – chop up program into parallel pieces
- Advantage 3: Modularity
 - More important than you might think
 - Chop large problem up into simpler pieces
 - » To compile, for instance, gcc calls cpp | cc1 | cc2 | as | ld
 - » Makes system easier to extend

1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.86

High-level Example: Web Server



- Server must handle many requests
- Non-cooperating version:

```
serverLoop() {
    con = AcceptCon();
    ProcessFork(ServiceWebPage(), con);
}
```
- What are some disadvantages of this technique?

1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.87

Threaded Web Server

- Now, use a single process
- Multithreaded (cooperating) version:

```
serverLoop() {
    connection = AcceptCon();
    ThreadFork(ServiceWebPage(), connection);
}
```
- Looks almost the same, but has many advantages:
 - Can share file caches kept in memory, results of CGI scripts, other things
 - Threads are *much* cheaper to create than processes, so this has a lower per-request overhead
- Question: would a user-level (say one-to-many) thread package make sense here?
 - When one request blocks on disk, all block...
- What about Denial of Service attacks or digg / Slash-dot effects?

1/30/20


Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.88



Thread Pools

- Problem with previous version: Unbounded Threads
 - When web-site becomes too popular – throughput sinks
- Instead, allocate a bounded “pool” of worker threads, representing the maximum level of multiprocessing



```

master() {
    allocThreads(worker, queue);
    while(TRUE) {
        con=AcceptCon();
        Enqueue(queue, con);
        wakeUp(queue);
    }
}


worker(queue) {
    while(TRUE) {
        con=Dequeue(queue);
        if (con==null)
            sleepOn(queue);
        else
            ServiceWebPage(con);
    }
}
    
```

1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.89

ATM Bank Server



- ATM server problem:
 - Service a set of requests
 - Do so without corrupting database
 - Don't hand out too much money

1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.90

ATM bank server example

- Suppose we wanted to implement a server process to handle requests from an ATM network:

```

BankServer() {
    while (TRUE) {
        ReceiveRequest(&op, &acctId, &amount);
        ProcessRequest(op, acctId, amount);
    }
}

ProcessRequest(op, acctId, amount) {
    if (op == deposit) Deposit(acctId, amount);
    else if ...
}

Deposit(acctId, amount) {
    acct = GetAccount(acctId); /* may use disk I/O */
    acct->balance += amount;
    StoreAccount(acct); /* Involves disk I/O */
}
    
```

- How could we speed this up?
 - More than one request being processed at once
 - Event driven (overlap computation and I/O)
 - Multiple threads (multi-proc, or overlap comp and I/O)

1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.91

Event Driven Version of ATM server

- Suppose we only had one CPU
 - Still like to overlap I/O with computation
 - Without threads, we would have to rewrite in event-driven style

- Example

```

BankServer() {
    while(TRUE) {
        event = WaitForNextEvent();
        if (event == ATMRequest)
            StartOnRequest();
        else if (event == AcctAvail)
            ContinueRequest();
        else if (event == AcctStored)
            FinishRequest();
    }
}
    
```

- What if we missed a blocking I/O step?
- What if we have to split code into hundreds of pieces which could be blocking?
- This technique is used for graphical programming

1/30/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 4.92

Can Threads Make This Easier?

- Threads yield overlapped I/O and computation without “deconstructing” code into non-blocking fragments
 - One thread per request
- Requests proceeds to completion, blocking as required:

```
Deposit(acctId, amount) {  
    acct = GetAccount(actId); /* May use disk I/O */  
    acct->balance += amount;  
    StoreAccount(acct);      /* Involves disk I/O */  
}
```

- Unfortunately, shared state can get corrupted:

<u>Thread 1</u>	<u>Thread 2</u>
load r1, acct->balance	
	load r1, acct->balance
	add r1, amount2
	store r1, acct->balance
add r1, amount1	
store r1, acct->balance	

Summary

- Processes have two parts
 - One or more Threads (Concurrency)
 - Address Spaces (Protection)
- Threads: unit of concurrent execution
 - Useful for parallelism, overlapping computation and IO, organizing sequences of interactions (protocols)
 - Require: multiple stacks per address space
 - Thread switch:
 - » Save/Restore registers, "return" from new thread's switch routine
- Concurrency accomplished by multiplexing CPU Time:
 - Unloading current thread (PC, registers)
 - Loading new thread (PC, registers)
 - Such context switching may be voluntary (yield(), I/O operations) or involuntary (timer, other interrupts)
- Concurrent threads introduce problems when accessing shared data
 - Programs must be insensitive to arbitrary interleavings
 - Without careful design, shared variables can become completely inconsistent