
CS 267: Introduction to Data Parallelism Lecture 6

Kathy Yelick

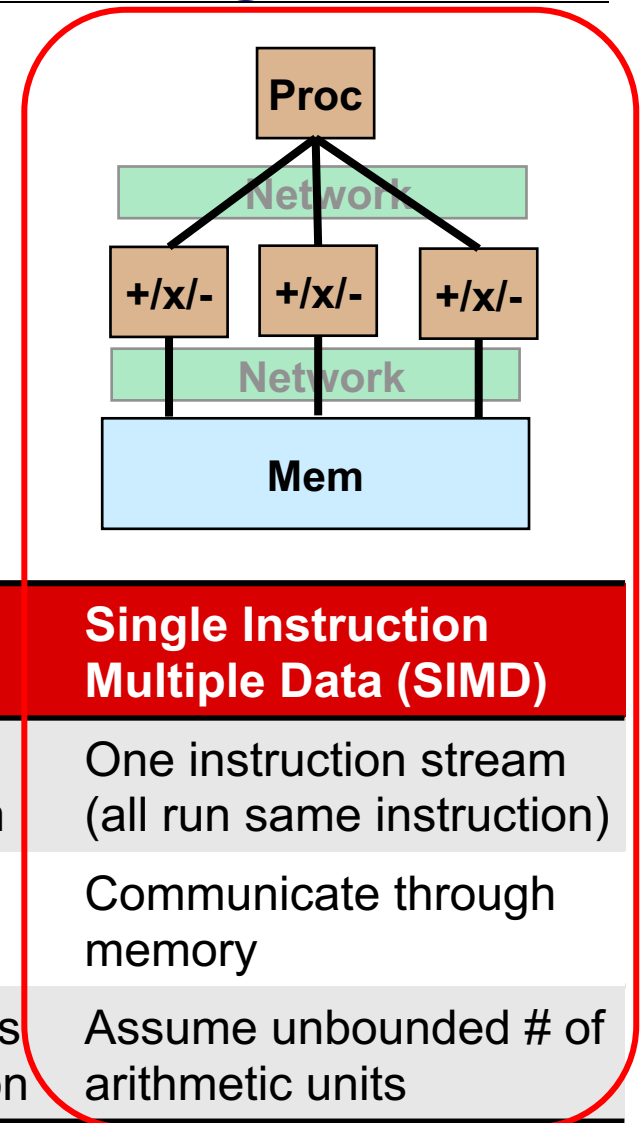


<https://sites.google.com/lbl.gov/cs267-spr2019/>

Lessons from Today's Lecture

- Data parallelism is beautiful!
- Automatically mapping it to (today's) hardware is hard
- Surprising things you can do with scans
- Many parallel programming models use some data parallel features
 - GPUs
 - MPI collectives
 - Cloud MapReduce
- Useful in designing (nontrivial) parallel algorithms

“Every nontrivial parallel algorithm uses a prefix scan”

Parallel Machines and Programming



Shared Memory

Processors execute own instruction stream

Communicate by reading/writing memory

Cost of a read/write is constant

Distributed Memory

Processors execute own instruction stream

Communicate by sending messages

Message time depends on size, but not location

Single Instruction Multiple Data (SIMD)

One instruction stream (all run same instruction)

Communicate through memory

Assume unbounded # of arithmetic units

- These are the natural “abstract” machine models

Data Parallel Programming: Unary Operators


- Unary operations applied to all elements of an array

A = array

B = array

f = square (any unary function, i.e., 1 argument)

B = f(A)



Data Parallel Programming: Binary Operators

- Binary operations applied to all pairs of elements


A = array

B = array

C = array

- or any other binary operator


$C = A - B$



Data Parallel Programming: Broadcast


- Broadcast fill a value into all elements of an array

a = scalar
B = a




- Useful for $a \cdot X + Y$ called axpy, saxpy, daxpy

$a \cdot X + Y$



Data Parallel Programming: Reduce


- Reduce an array to a value with + or any associative op



- Associative so we can perform op in any order
- Useful for dot products (ddot, sdot, etc.)


$$\mathbf{b} = \mathbf{X}^T \mathbf{Y} = \sum_j \mathbf{X}[j] * \mathbf{Y}[j]$$

$$\mathbf{b} = \text{dot}(\mathbf{X}, \mathbf{Y}) = \text{sum}(\mathbf{X} .* \mathbf{Y})$$




Data Parallel Programming: Scans

- Fill array with partial reductions any associative op
- Sum scan:



- Max scan:



Inclusive and Exclusive Scans

Two variations of a scan, given an input vector $[x_0, x_1, \dots, x_{n-1}]$:

- **inclusive** scan includes input x_i when computing output y_i

$$[a_0, (a_0 \odot a_1), \dots, (a_0 \odot a_1 \dots \odot a_{n-1})]$$

e.g., `add_scan_inclusive([1, 0, 3, 0, 2])` \rightarrow `[1, 1, 4, 4, 6]`

- **exclusive** scan does *not* x_i when computing output y_i

$$[I, a_0, (a_0 \odot a_1), \dots, (a_0 \odot a_1 \dots \odot a_{n-2})] \text{ where } I \text{ is the identity for } \odot$$

e.g., `add_scan_exclusive([1, 0, 3, 0, 2])` \rightarrow `[0, 1, 1, 4, 4]`

You can get the inclusive version from the exclusive by applying the operation across vectors: $X \odot Y$. You can convert both directions using vector shifts left or right.

Data Parallel Programming: Masks

- Can apply operations under a “mask”

M = array of 0/1 (True/False)

A = array

B = array

A = A + B under M

A:

3	1	1	2	3	3	4	2	2	2	1	3	1	1	1	3	3	2	1
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

B:

0	1	1	4	1	0	2	1	4	3	1	0	1	1	2	3	5	3	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

M:

1	0	0	1	1	0	0	0	1	1	1	1	0	0	1	0	0	1	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---



+ under mask

A:


0			4	1				4	3	1	0			2			3	
3	1	1	6	4	3	4	2	6	5	2	3	1	1	3	3	3	5	1

- Related: Segmented scans to be presented later

Idealized Hardware and Performance Model


SIMD Systems Implemented Data Parallelism

- A large number of (usually) tiny processors.
 - **A single “control processor” issues each instruction.**
 - **Each processor executes the same instruction.**
 - **Some processors may be turned off on some instructions.**
- Originally machines were specialized to scientific computing, few made (CM2, Maspar)



Ideal Cost Model for Data Parallelism


- Machine
 - An unbounded number of processors (p)
 - Control overhead is free
 - Communication is free
- Shows the inherent parallelism (inherent serialization)
- Called the algorithm's "span"
- Defines a lower bound on real machines



Cost on Ideal Machine (Span)

- Span for unary or binary operations (pleasingly parallel)


$$C = A + B$$



Cost $O(1)$
since p is unbounded

- Even if arrays are not aligned, communication is “free” here
- Reductions and broadcasts

$$s = \text{sum}(C)$$




Cost $O(\log(n))$


Using a tree of
processors

Broadcast and reduction use processor trees

- **Broadcast** of 1 value to p processors with $\log n$ span




- **Reduction** of n values to 1 with $\log n$ span
- Takes advantage of associativity in $+$, $*$, \min , \max , etc.




Can reductions go faster: No, $\log n$ lower bound and on any function of n variables!

- Given a function $f(x_1, \dots, x_n)$ of n input variables and 1 output variable, how fast can we evaluate it in parallel?
- Assume we only have binary operations, one per time step
- After 1 time step, an output can only depend on two inputs
- Use induction to show that after k time units, an output can only depend on 2^k inputs
 - After $\log_2 n$ time units, output depends on at most n inputs
- A binary tree performs such a computation




Multiplying n-by-n matrices in $O(\log n)$ time

- Use n^3 processors
- Step 1: For all $(1 \leq i, j, k \leq n)$ $P(i, j, k) = A(i, k) * B(k, j)$
 - cost = 1 time unit, using n^3 processors
- Step 2: For all $(1 \leq i, j \leq n)$ $C(i, j) = \sum_{k=1}^n P(i, j, k)$
 - cost = $O(\log n)$ time, using n^2 trees, $n^3 / 2$ processors each



**Put a processor
at every point in
this cube**

Related to Communication-Optimal “2.5D” MatMul



- Processors execute internal sub-cubes

What about Scan (aka Parallel Prefix)?

- Recall: the **scan** operation takes a **binary associative** operator \odot , and an array of n elements

$$[a_0, a_1, a_2, \dots a_{n-1}]$$

and produces the array

$$[a_0, (a_0 \odot a_1), \dots (a_0 \odot a_1 \dots \odot a_{n-1})]$$

- Example: **add scan** of

$$[1, 2, 0, 4, 2, 1, 1, 3] \quad \text{is} \quad [1, 3, 3, 7, 9, 10, 11, 14]$$

- Other operators
 - Reals: +, *, min, max
 - Booleans: and, or
 - Matrices: mat mul

Can we parallelize a scan?

- It looks like this:

```
y(0) = 0;  
for i = 1:n  
    y(i) = y(i-1) + x(i);
```

- Takes $n-1$ operations (adds) to do in serial
- The i^{th} iteration of the loop depends completely on the $(i-1)^{\text{st}}$ iteration.
- Impossible to parallelize, right?

A clue

input = (1, 2, 3, 4, 5, 6, 7, 8)

output = (1, 3, 6, 10, 15, 21, 28, 36)


Is there any value in adding, say, $5+6+7+8$?

If we separately have $1+2+3+4$, what can we do?


Suppose we added $1+2$, $3+4$, etc. pairwise -- what could we do?

Prefix sum in parallel

Algorithm: 1. Pairwise sum 2. Recursive prefix 3. Pairwise sum



Parallel prefix cost



Pairwise sum

Recursive prefix

**Pairwise sum
(update odds)**

Time for this algorithm on one processor (work)

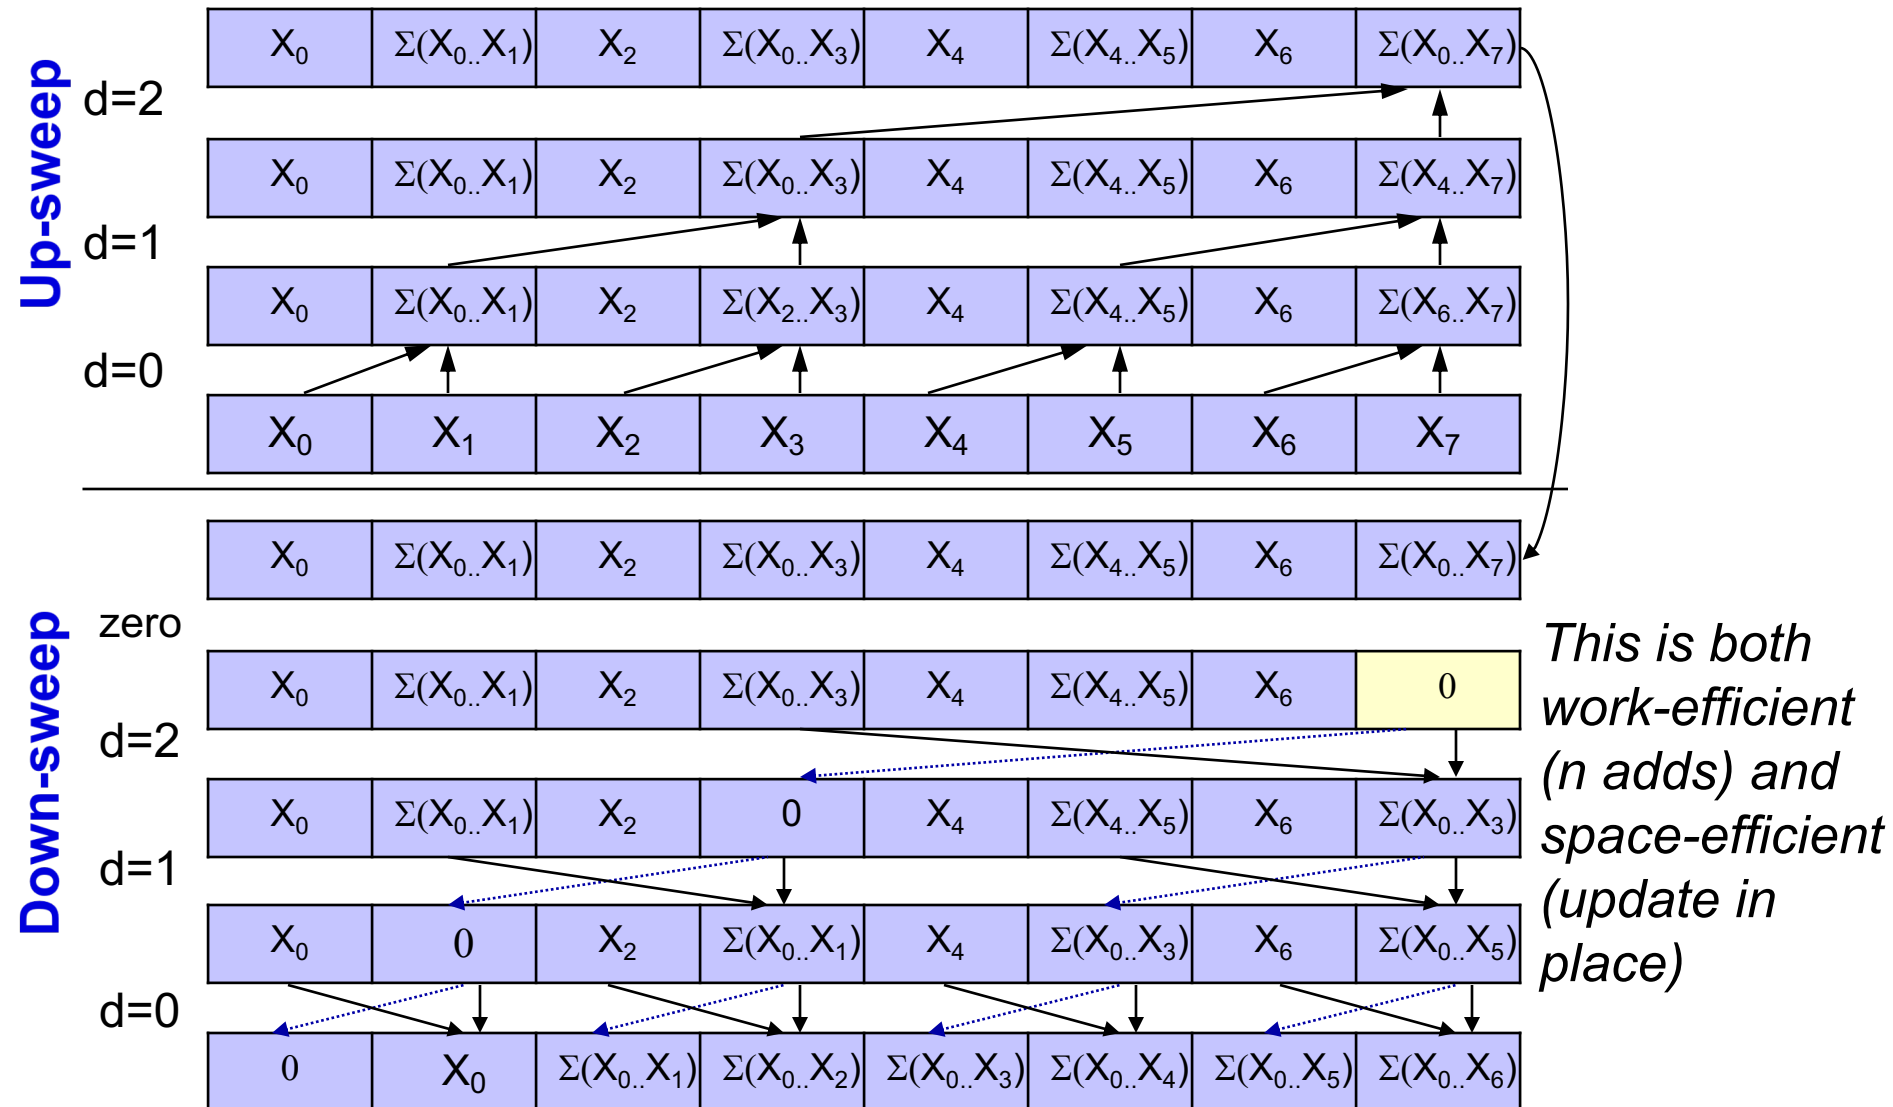
- $T_1(n) = n/2 + n/2 + T_1(n/2) = n + T_1(n/2) = 2n - 1$

Time on unbounded number of processors (span)

- $T_\infty(n) = 2 \log n$

Parallelism at the cost of more work (2x)!


Non-recursive view of parallel prefix scan




Real Hardware (Today)

Vectors use Data Parallelism (at smaller scale)

- Vector instructions operate on a vector of elements
 - These are specified as operations on vector registers




- Old supercomputer vector register: ~32-64 elts
 - The number of elements is larger than the amount of parallel hardware, called vector **pipes** or **lanes**, say 2-4
- The hardware performs a full vector operation in
 - #elements-per-vector-register / #pipes
 - E.g., 64 elements in register, but only 8 fp adders to use
 - “Virtualizes” the amount of hardware, which is n



Cray X1: *Parallel Vector Architecture*

Cray combined several technologies in the X1

- **12.8 Gflop/s Vector processors (MSP)**
- **Shared caches (unusual on earlier vector machines)**
- **4 processor nodes sharing up to 64 GB of memory**
- **Single System Image to 4096 Processors**
- **Remote put/get between nodes (faster than MPI)**




- **Expensive to design and build, market too small**

Vectors use Data Parallelism (at smaller scale)

- **SIMD instructions on microprocessors are vectors**
 - Shorter than old vector supercomputers (e.g., 256 bits)
 - They don't virtualize the hardware (arithmetic units), so each processor version may require code rewrites
- **Reductions and broadcasts are in register**
 - Require inside-register data movement
- **Assuming vector length (or SIMD width) are small constants → no theoretical speedup**
 - But in practice this can make a big different (2-16x...)
 - And algorithms may still be useful
- **Revisit these ideas with GPUs**

Data parallelism on Distributed Memory



Distributed Memory

Processors execute
own instruction stream

Communicate by
sending messages

Message time depends
on size, but not location

- **Today's parallel machines**
 - Powerful processors
 - Distributed memory (at scale)
 - Clusters or MPPs (Massively Parallel Processors)
- **Need to map n-way parallelism to p-way**
 - Attempts to do this automatically
- **High Performance Fortran**
 - Large effort in the 90s
 - Semi-automatic: Data layout hints were necessary
 - And it was still hard
- **But still useful manually**

Mapping Data Parallelism to Clusters

- Binary and unary operations on MPPs

$$C = A + B$$

A: 

+

B: 

=

C: 

Cost $O(n/p)$

p speedup

- If arrays are not “aligned” then communication required
- Reductions and broadcasts

C: 

Cost $O(n/p)$



local sum

+ $\log p$)


$$s = \text{sum}(C)$$

s: 

tree reduction

Almost p speedup

Parallel prefix cost on p processors



**Compute local
prefix sums in
 n/p steps**

**Updates across
processors in
 $\log p$ steps**

Time for this algorithm in parallel:

- $T_p(n) = O(n/p + \log p)$

**serial time
on each
processor**

**communication and
computation up and
down the processor tree**

The myth of $\log n$

- The $\log_2 n$ span is **not** the main reason for the usefulness of parallel prefix.
- **Say $n = 1,000,000p$** (1,000,000 elements per proc)
 - Cost = (2000000 adds) + ($\log_2 P$ message passings)

↑

fast & embarrassingly parallel

(2000000 local adds are serial for each processor, of course)

Key to implementing data parallel algorithms on clusters, SMPs, MPPs, i.e., modern supercomputers

Data Parallelism is an Elegant Programming Model

- Strict data parallelism has serial semantics:
 - E.g., no difference from executing $A+B$ one element at a time or in parallel
- Reductions also preserve serial semantics for truly associative operations:
 - $+$ $*$ \min , etc. on integers and more;
 - some differences for floating point due to order of evaluation (but can be deterministic, i.e., the same result every time)
- Easy to understand and reason about
- “In spirit” in MPI collectives, CUDA, MapReduce...

Limitations:

- Some algorithms (e.g., adaptive) don't fit easily
- Non-trivial to implement on some hardware

Scans are useful for many things (partial list here)

- Reduction and broadcast in $O(\log n)$ time
- Parallel prefix (scan) in $O(\log n)$ time
- Adding two n -bit integers in $O(\log n)$ time
- Multiplying n -by- n matrices in $O(\log n)$ time
- Inverting n -by- n triangular matrices in $O(\log^2 n)$ time
- Inverting n -by- n dense matrices in $O(\log^2 n)$ time
- Evaluating arbitrary expressions in $O(\log n)$ time
- Evaluating recurrences in $O(\log n)$ time
- “2D parallel prefix”, for image segmentation (Catanzaro & Keutzer)
- Sparse-Matrix-Vector-Multiply (SpMV) using Segmented Scan
- Parallel page layout in a browser (Leo Meyerovich, Ras Bodik)
- Solving n -by- n tridiagonal matrices in $O(\log n)$ time
- Traversing linked lists
- Computing minimal spanning trees
- Computing convex hulls of point sets...

Application: Stream Compression

- Given an array of 0/1 flags

flags =

1	0	1	1	0	0	1	1
---	---	---	---	---	---	---	---

and an array (stream) of values

values =

3	2	4	1	5	3	3	1
---	---	---	---	---	---	---	---

compress into

result =

3	4	1	3	1
---	---	---	---	---


- Step 1: Compute an exclusive add scan of flags:

index =

0	1	1	2	3	3	3	4
---	---	---	---	---	---	---	---

- Step 2: “Scatter” values into result at index, masked by flags

result[index] = values at flags



Application: Radix Sort

4	7	2	6	3	5	1	0
---	---	---	---	---	---	---	---

input

0	1	0	0	1	1	1	0
---	---	---	---	---	---	---	---

odds = last bit of each element

1	0	1	1	0	0	0	1
---	---	---	---	---	---	---	---

evens = complement of odds (last bit = 0)

0	1	1	2	3	3	3	3
---	---	---	---	---	---	---	---

epos = exclusive sum scans of evens

4	4	4	4	4	4	4	4
---	---	---	---	---	---	---	---

totalEvens = broadcast last element

0	1	2	3	4	5	6	7
---	---	---	---	---	---	---	---

indx = constant array of 0..n

0-0+4 =4	1-1+4 =4	2-1+4 =5	3-2+4 =5	4-3+4 =5	5-3+4 =6	6-3+4 =7	7-4+4 =7
-------------	-------------	-------------	-------------	-------------	-------------	-------------	-------------

4	4	5	5	5	6	7	7
---	---	---	---	---	---	---	---


oddpos = indx – epos + totalEvens

0	4	1	2	5	6	7	3
---	---	---	---	---	---	---	---

pos = if evens then esum else oddpos

4	7	2	6	3	5	1	0
---	---	---	---	---	---	---	---

*Using two
masked
assignments*



Scatter input using pos as index

Repeat with next bit to left until done

Application: Radix Sort

4	7	2	6	3	5	1	0
---	---	---	---	---	---	---	---

Sort on least significant bit

4	2	6	0	7	3	5	1
---	---	---	---	---	---	---	---

$XX0 < XX1$ (evens before odds)

Bit₂=0 Bit₂=1

4	2	6	0	7	3	5	1
---	---	---	---	---	---	---	---

Sort on next bit $X0X < X1X$

4	0	5	1	2	6	7	3
---	---	---	---	---	---	---	---

Bit₁=0 Bit₁=1

4	0	5	1	2	6	7	3
---	---	---	---	---	---	---	---

Sort on next bit

0	1	2	3	4	5	6	7
---	---	---	---	---	---	---	---

$0XX < 1XX$ (<4 before ≥ 4)

Bit₀=0 Bit₀=1

*Each step maintains the ordering
unless, they have to switch based on
the current bit*

Application: Adding n-bit integers in $O(\log n)$ time

- Computing sum s of two n -bit binary numbers, a and b
 - $a = a[n-1] a[n-2] \dots a[0]$ and $b = b[n-1] b[n-2] \dots b[0]$
 - $s = a+b = s[n] s[n-1] \dots s[0]$ (use carry-bit array $c = c[n-1] \dots c[0] c[-1]$)

- Formula

$c[-1] = 0$... rightmost carry bit

for $i = 0$ to $n-1$

$s[i] = (a[i] \text{ xor } b[i]) \text{ xor } c[i-1]$

$c[i] = ((a[i] \text{ xor } b[i]) \text{ and } c[i-1]) \text{ or } (a[i] \text{ and } b[i])$... next carry bit

- Example

- $a = 22$

- $b = 29$

$a = 1\ 0\ 1\ 1\ 0$ (22)

$b = 1\ 1\ 1\ 0\ 1$ (29)

$c = 1\ 1\ 1\ 0\ 0\ 0\ 0$

$s = 1\ 1\ 0\ 0\ 1\ 1$ (51)

- Challenge: compute all $c[i]$ in $O(\log n)$ time via parallel prefix

Application: Adding n-bit integers in $O(\log n)$ time

- Recall carry bit calculation

$c[-1] = 0$... rightmost carry bit

for $i = 0$ to $n-1$

$c[i] = ((a[i] \text{ xor } b[i]) \text{ and } c[i-1]) \text{ or } (a[i] \text{ and } b[i])$... next carry bit

- Compute all $c[i]$ in $O(\log n)$ time via parallel prefix

for all $(0 \leq i \leq n-1)$ $p[i] = a[i] \text{ xor } b[i]$... propagate bit

for all $(0 \leq i \leq n-1)$ $g[i] = a[i] \text{ and } b[i]$... generate bit

$$\begin{bmatrix} c[i] \\ 1 \end{bmatrix} = \begin{bmatrix} (p[i] \text{ and } c[i-1]) \text{ or } g[i] \\ 1 \end{bmatrix} = \begin{bmatrix} p[i] & g[i] \\ 0 & 1 \end{bmatrix} * \begin{bmatrix} c[i-1] \\ 1 \end{bmatrix} = M[i] * \begin{bmatrix} c[i-1] \\ 1 \end{bmatrix}$$


$$= M[i] * M[i-1] * \dots * M[0] * \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

... evaluate $M[i] * M[i-1] * \dots * M[0]$ by parallel prefix

... 2-by-2 Boolean matrix multiplication is associative

- Used in all computers to -- Carry look-ahead addition

This idea is used in all hardware



- Even going back to Babbage

Segmented Scans

Inputs = value array, flag array,
associative operator \oplus

Inclusive segmented sum scan

1	2	3	4	5	6	7	8
0	0	1	0	0	1	0	1

**Flags are sometimes done with
Boolean and switch points**

F	F	T	T	T	F	F	T
---	---	---	---	---	---	---	---

Result

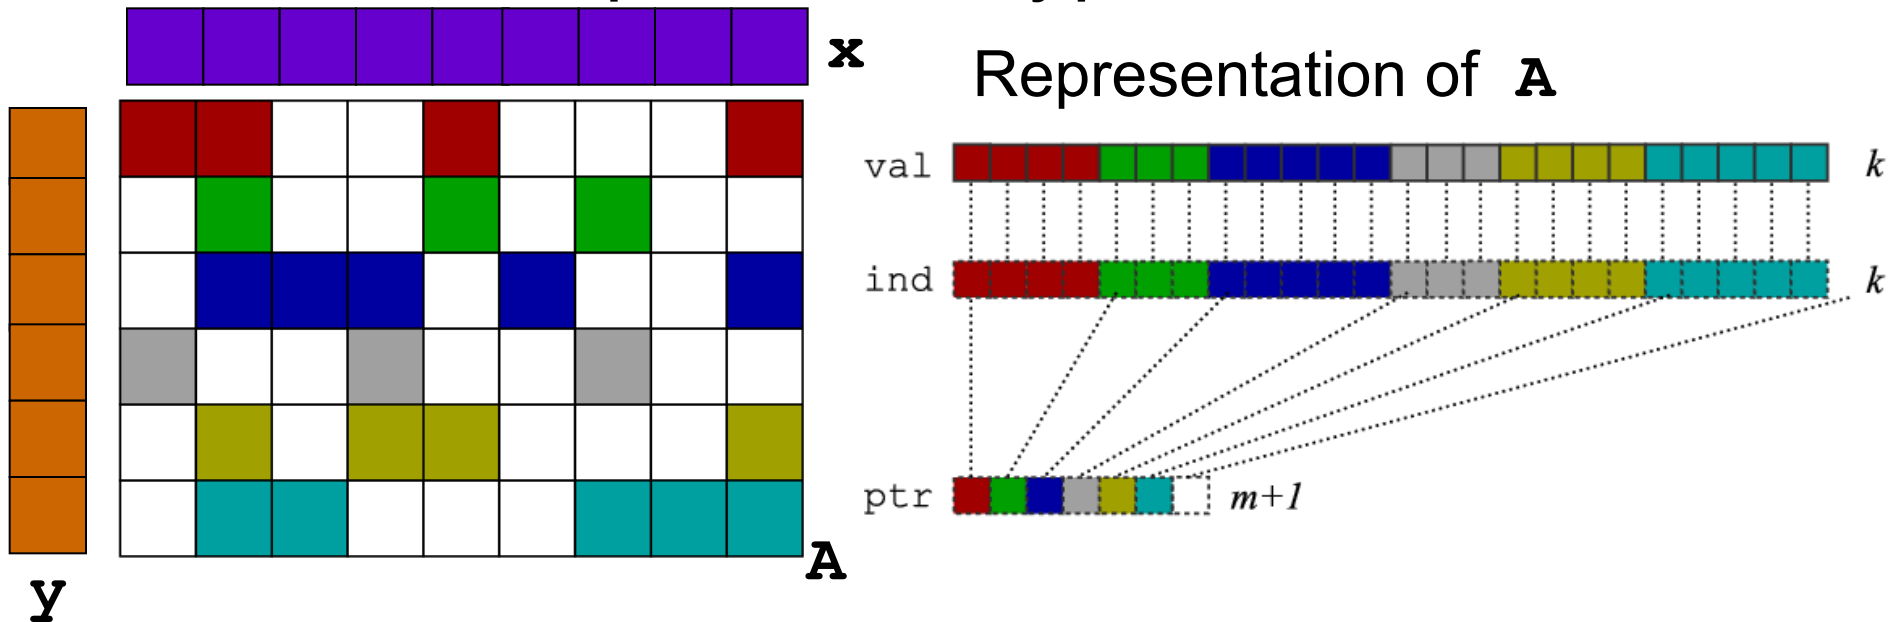
1	3	3	7	12	6	13	8
---	---	---	---	----	---	----	---

SpMV in Compressed Sparse Row (CSR) Format

SpMV: $y = y + A \cdot x$

Sparse matrices: only store, do arithmetic, on nonzero entries

CSR format is simplest one of many possible data structures for A



Matrix-vector multiply kernel: $y(i) \leftarrow y(i) + A(i,j) \times x(j)$

for each row i

for $k=ptr[i]$ to $ptr[i+1]-1$ do

$y[i] = y[i] + val[k] * x[ind[k]]$

SPMV (Segmented Suffix Scan)

- Sparse Matrix-Vector Multiplication (SPMV)

x=

0	1	2	3	4	5	6
1	2	1	2	1	2	1

PTR =

0	2	5	7	9
---	---	---	---	---

Row Starts

IND =

1	2	0	3	4	1	2	5	6	0	1	4
---	---	---	---	---	---	---	---	---	---	---	---

VAL =

1	1	1	2	2	2	1	2	2	1	3	1
---	---	---	---	---	---	---	---	---	---	---	---

PROD =

1	1	1	2	2	2	2	2				
---	---	---	---	---	---	---	---	--	--	--	--

FLAG = ZEROS + ONES (PTR)

PROD = VAL * X (IND)

**SUMS = SUM_SUFFIX (PROD, SEGMENT=SEGS
PREFIX (FLAG))**

Y = SUMS (PTR)

Segmented Operations for Sparse Matrix Computation on Vector Multiprocessors . Guy E. Blelloch, Michael A. Heroux, and Marco Zagha. CMU-CS-93-173

Application: Fibonacci via Matrix Multiply Prefix

$$\boxed{F_{n+1} = F_n + F_{n-1}}$$

$$\boxed{\begin{pmatrix} F_{n+1} \\ F_n \end{pmatrix} = \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix} \begin{pmatrix} F_n \\ F_{n-1} \end{pmatrix}}$$

Can compute all F_n by matmul_prefix on

$$\left[\begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}, \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}, \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}, \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}, \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}, \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}, \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}, \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}, \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix} \right]$$

then select the upper left entry

Slide source: Alan Edelman

Lexical analysis (tokenizing, scanning)

- Given a language of:
 - Identifiers: string of chars
 - Strings: in double quotes
 - Ops: +, -, *, =, <, >, <=, >=

TABLE I. A Finite-State Automaton for Recognizing Tokens

Old State	Character Read													New line
	A	B	...	Y	Z	+	-	*	<	>	=	"	Space	
N	A	A	...	A	A	*	*	*	<	<	*	Q	N	N
A	Z	Z	...	Z	Z	*	*	*	<	<	*	Q	N	N
Z	Z	Z	...	Z	Z	*	*	*	<	<	*	Q	N	N
*	A	A	...	A	A	*	*	*	<	<	*	Q	N	N
<	A	A	...	A	A	*	*	*	<	<	=	Q	N	N
=	A	A	...	A	A	*	*	*	<	<	*	Q	N	N
Q	S	S	...	S	S	S	S	S	S	S	S	E	S	S
S	S	S	...	S	S	S	S	S	S	S	S	E	S	S
E	E	E	...	E	E	*	*	*	<	<	*	S	N	N

- Lexical analysis
 - Replace every character in the string with the array representation of its state-to-state function (column).
 - Perform a parallel-prefix operation with \oplus as the array composition. Each character becomes an array representing the state-to-state function for that prefix.
 - Use initial state (row 1) to index into these arrays.

Hillis and Steele, CACM 1986

Inverting Dense n-by-n matrices in $O(\log^2 n)$ time

- Lemma 1: Cayley-Hamilton Theorem
 - expression for A^{-1} via characteristic polynomial in A
- Lemma 2: Newton's Identities
 - Triangular system of equations for coefficients of characteristic polynomial, where matrix entries = s_k
- Lemma 3: $s_k = \text{trace}(A^k) = \sum_{i=1}^n A^k[i,i]$
- Csanky's Algorithm (1976)
 - 1) Compute the powers A^2, A^3, \dots, A^{n-1} by parallel prefix
cost = $O(\log^2 n)$
 - 2) Compute the traces $s_k = \text{trace}(A^k)$
cost = $O(\log n)$
 - 3) Solve Newton identities for coefficients of characteristic polynomial
cost = $O(\log^2 n)$
 - 4) Evaluate A^{-1} using Cayley-Hamilton Theorem
cost = $O(\log n)$

○ **Completely numerically unstable**

Lessons from Data Parallel Languages

- Sequential semantics (or nearly) is very nice
 - Debugging is much easier without non-determinism
 - Correctness easier to reason about
- Cost model is independent of number of processors
 - How much inherent parallelism
- Need to “throttle” parallelism
 - $n \gg p$ can be hard to map, especially with nesting
 - Memory use is a problem

See: Blelloch “NESL Revisited”, Intel Workshop 2006