

---

**CS 267**  
**Lecture 17b: Parallel Machine**  
**Learning, Part 2**  
**(Unsupervised Learning)**

**Aydin Buluc**

**<https://sites.google.com/lbl.gov/cs267-spr2019/>**

## Outline of the lecture

---

### **Tuesday: Part 1, Intro and Supervised Learning**

- Machine Learning & Parallelism Intro
- Neural Network Basics
- Scaling Deep Neural Network Training
- Support Vector Machines


### **Today: Part 2, Unsupervised Learning**

- Non-Negative Matrix Factorization
- Spectral and Markov Clustering
- Sparse Inverse Covariance Matrix Estimation

# Non-negative matrix factorization (NMF)

---

$$\min_{W \geq 0, H \geq 0} f(W, H) = \frac{1}{2} \|A - WH\|_F^2$$



- Dimensionality reduction with non-negativity constraints
- The name “factorization” is a misnomer; NMF is just a low-rank approximation as exact factorization is NP-hard
- NMF is a family of methods, not just one algorithm

# The Alternating Updates Framework

---

**Initialize H**

**Repeat until convergence:**

1. For fixed H, solve  $\min_{W \geq 0} \|H^T W^T - A^T\|_F^2$
2. For fixed W, solve  $\min_{H \geq 0} \|WH - A\|_F^2$

Lots of algorithms fall into this framework.

- Multiplicative update (MU)
- Alternating least squares (ALS)
- Alternating non-negative least squares (ANLS)

J. Kim and H. Park. "Fast nonnegative matrix factorization: An active-set-like method and comparisons." SIAM Journal on Scientific Computing, 2011

**Caveat emptor:** This is not the only method for finding an NMF

Gemulla, Rainer, et al. "Large-scale matrix factorization with distributed stochastic gradient descent." KDD, 2011

# The Alternating Updates Framework

Main computation is large-scale matrix multiplications:

1.  $\mathbf{H}\mathbf{H}^T$  and  $\mathbf{A}\mathbf{H}^T$  for updating  $\mathbf{W}$ , given a fixed  $\mathbf{H}$
  2.  $\mathbf{W}^T\mathbf{W}$  and  $\mathbf{W}^T\mathbf{A}$  for updating  $\mathbf{H}$ , given a fixed  $\mathbf{W}$
- In general  $\mathbf{W}$  and  $\mathbf{H}$  are dense, but short-fat or tall-skinny
  - $\mathbf{A}$  is often sparse but can be dense depending on application
  - For increased interpretability,  $\mathbf{H}$  or  $\mathbf{W}$  can also be sparse




Figure: Kim and Park


# The Alternating Updates Framework

Main computation is large-scale matrix multiplications


Choose the best distribution and algorithm depending on:

- 1- the relative sizes of the dimensions of the matrices
- 2- the number of processors

**This work:** Never communicate  $A$ , because it is asymptotically larger than  $H$  and  $W$



(a) 1D Distribution with  $p = p_r = 4$  and  $p_c = 1$ .



(b) 2D Distribution with  $p_r = 3$  and  $p_c = 2$ .

Kannan, Ballard, Park. "MPI-FAUN: An MPI-Based Framework for Alternating-Updating Nonnegative Matrix Factorization". 2016.

## Outline of the lecture

---

### **Tuesday: Part 1, Intro and Supervised Learning**

- Machine Learning & Parallelism Intro
- Neural Network Basics
- Scaling Deep Neural Network Training
- Support Vector Machines

### **Today: Part 2, Unsupervised Learning**

- Non-Negative Matrix Factorization
- Spectral and Markov Clustering
- Sparse Inverse Covariance Matrix Estimation

## Many families of methods


- Centroid based (k-means, k-medians, and variations)
  - Flow based (Markov clustering)
  - Spectral methods
  - Density based (DBSCAN, OPTICS)
  - Agglomerative methods (single linkage clustering)
  - ...
- 
- Often the right method depends on the input characteristics and require some domain knowledge.
  - We will talk about parallel algorithms for two: **Spectral Clustering** and **Markov Clustering (MCL)**.



## Spectral Clustering

---

- **Input:** Similarities between data points
- Many ways to compute similarity, some are domain specific: cosine, Jaccard index, Pearson correlation, Spearman's rho, Bhattacharyya distance, LOD score, ...
- We can represent the relationships between data points in a graph.
- Weight the edges by the similarity between points



## Graph definitions

---

- $\varepsilon$ -neighborhood graph
  - Identify a threshold value,  $\varepsilon$ , and include edges if the affinity between two points is greater than  $\varepsilon$ .
- k-nearest neighbors
  - Insert edges between a node and its k-nearest neighbors.
  - Each node will be connected to (at least) k nodes.
- Fully connected
  - Insert an edge between every pair of nodes.

## Spectral Clustering Intuition

---

- The minimum cut of a graph identifies an optimal partitioning of the data.
- Spectral Clustering
  - Recursively partition the data set
    - Identify the minimum cut
    - Remove edges
    - Repeat until k clusters are identified
- **Problem:** Identifying a minimum cut is NP-hard.
- There are efficient approximations using linear algebra, based on the Laplacian Matrix, or **graph Laplacian**

# The Graph Laplacian

## □ Graph Laplacian


■ unnormalized graph Laplacian :  $L = D - W$

■ normalized graph Laplacian

$$L_{sym} = D^{-1/2} L D^{-1/2} = I - D^{-1/2} W D^{-1/2}$$

$$L_{rw} = D^{-1} L = I - D^{-1} W \quad \longleftarrow \text{related to random walk}$$

■ Example



$$L = \begin{pmatrix} 2 & -1 & -1 & 0 & 0 & 0 \\ -1 & 4 & -1 & -1 & -1 & 0 \\ -1 & -1 & 4 & 0 & -1 & -1 \\ 0 & -1 & 0 & 1 & 0 & 0 \\ 0 & -1 & -1 & 0 & 3 & -1 \\ 0 & 0 & -1 & 0 & -1 & 2 \end{pmatrix}$$

Assume the weights of edges are 1.

# One Spectral Clustering Algorithm

---

□ I recommend the **normalized symmetric Laplacian**, as the numerical eigenvalue problem there is easier to solve.

□ Normalized Spectral Clustering [Ng2002]


1. Construct a similarity graph and compute the normalized graph Laplacian  $L_{sym}$ .
2. Compute the  $k$  smallest eigenvectors  $u_1, u_2, \dots, u_k$  of  $L_{sym}$ .
3. Let  $U = [u_1 \ u_2 \ \dots \ u_k] \in \mathbb{R}^{n \times k}$ .
4. Normalized the rows of  $U$  to norm 1.

$$U_{ij} \leftarrow \frac{U_{ij}}{(\sum_k U_{ik}^2)^{1/2}}$$

5. Let  $y_i \in \mathbb{R}^k$  be the vector corresponding to the  $i$ th row of  $U$ .
6. Thinking of  $y_i$ 's as points in  $\mathbb{R}^k$ , cluster them with  $k$ -means algorithms.

## Why does it work? One intuitive explanation

- Ideal Case



$$L = D - W$$

2	-1	-1	0	0	0
-1	2	-1	0	0	0
-1	-1	2	0	0	0
0	0	0	2	-1	-1
0	0	0	-1	2	-1
0	0	0	-1	-1	2

$$Lv = \lambda v$$

1	0
1	0
1	0
0	1
0	1
0	1

- The multiplicity of the eigenvalue 0 gives the number of clusters (in this ideal case: the number of connected components).
- The real case is assumed to be an approximation to this situation.

## How to compute those smallest Eigenvectors?

---

- Implementation via the Lanczos Algorithm
  - Workhorse is sparse-matrix-vector (SpMV) multiply
  - SpMV has no/minimal data reuse, bound by communication
  - To optimize sparse-matrix-vector multiply and minimize its communication, we graph partition (next lecture)
- Alternative algorithms are possible
  - Power iteration is cheaper but numerically unstable
  - LOBPCG (Locally-Optimized Block Preconditioned Conjugate Gradient) uses sparse-matrix times multiple vectors, thus has more favorable performance profile due to possible data reuse.
- In the end, you probably just want to call something existing.
  - ARPACK implements reverse communication eigensolvers: You implement the SpMV, it implements the numerical outer logic
  - PARPACK is its parallel version, The following code uses it:  
<https://github.com/openbigdatagroup/pspectralclustering>

## Outline of the lecture

---

### **Tuesday: Part 1, Intro and Supervised Learning**

- Machine Learning & Parallelism Intro
- Neural Network Basics
- Scaling Deep Neural Network Training
- Support Vector Machines


### **Today: Part 2, Unsupervised Learning**

- Non-Negative Matrix Factorization
- Spectral and Markov Clustering
- Sparse Inverse Covariance Matrix Estimation



# Philosophy of the Markov Cluster Algorithm (MCL)

---



The number of **edges or higher-length paths** between two arbitrary nodes in a cluster is greater than the number of paths between nodes from different clusters



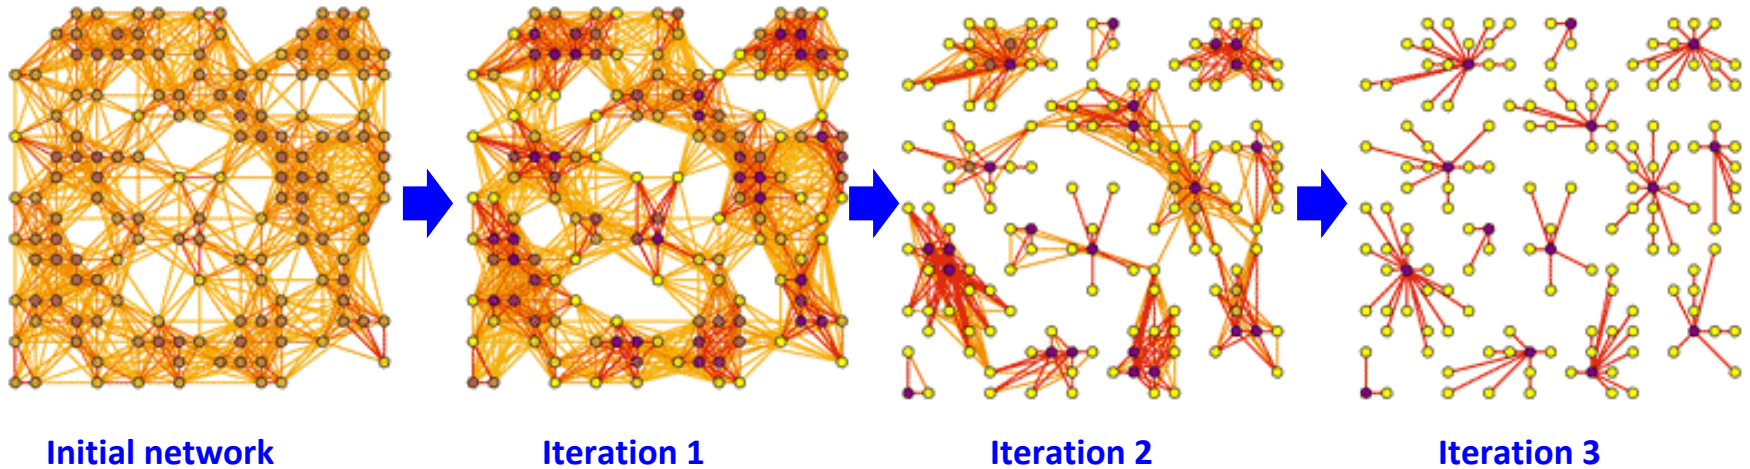
**Random walks** on the graph will frequently remain within a cluster



The algorithm **computes the probability** of random walks through the graph and **removes lower probability terms** to form clusters.

# The MCL Algorithm

Input: Adjacency matrix  $A$  (sparse & **column stochastic**)



**At each iteration:**

**Step 1 (Expansion):** Squaring the matrix

[corresponds to computing random walks of higher length]

**Step 2 (Inflation) :** Hadamard power of a matrix (taking powers entrywise)


[boost the probabilities of intra-cluster walks and demote inter-cluster walks]

## The expansion step of the MCL algorithm

---


- ❑ Goal: Compute random walks of higher length
- ❑ Input: A column stochastic matrix (A)
- ❑ Algorithm
  1. Sparse matrix-sparse matrix multiplication (**SpGEMM**):  $A^2$
  2. **Sparsify**  $A^2$  by removing low probability terms
    - **Prune** entries in  $A^2$  that are smaller than a threshold
    - **Recover** (if overdone pruning): Keep at least R entries (column-wise **top-K selection** )
    - **Selection** (if underdone pruning): Sparsify denser columns by keeping at most S entries (column-wise **top-K selection** )
- ❑ After sparsification at most  $\max(R, S)$  (**default to 1400**) entries remains in each column of  $A^2$

## A combined expansion and pruning step



- $b$ : number of columns in the output constructed at once
  - Smaller  $b$ : less parallelism, memory efficient ( $b=1$  is equivalent to sparse matrix-sparse vector multiplication used in MCL)
  - Larger  $b$ : more parallelism, memory intensive


## A combined expansion and pruning step



- $b$ : number of columns in the output constructed at once
  - HipMCL selects  $b$  dynamically as permitted by the available memory
  - The algorithm works in  $h=N/b$  phases where  $N$  is the number of columns (vertices in the network) in the matrix

# Current sparse matrix-matrix multiply algorithm in HipMCL

- ❑ Sparse SUMMA algorithm.
- ❑ Do this for each phase.
- ❑ Issue: repeated broadcast of A.
- ❑ Better ideas are brewing.




## Other algorithmic steps of HipMCL

---

- ❑ There is more than sparse matrix multiply here.
  - Parallel k-selection algorithm for each column of the matrix (for Recovery and Selection). **We will cover some algorithms in the sorting and searching lecture**
  - Parallel pruning algorithm
  - Parallel connected component algorithm (to identify clusters after MCL is converged). **Very fundamental graph algorithm, though we won't cover it this year.**
  - Parallel file I/O. **Reading terabytes of text efficiently is a challenge.**

**Azad, A.**, Pavlopoulos, G.A., Ouzounis, C.A., Kyrpides, N.C. and **Buluç, A.**, 2018. HipMCL: a high-performance parallel implementation of the Markov clustering algorithm for large-scale networks. Nucleic acids research.

# Scaling of HipMCL



On 1024 nodes, we need fewer phases because of more aggregated memory



## Outline of the lecture

---

### **Tuesday: Part 1, Intro and Supervised Learning**

- Machine Learning & Parallelism Intro
- Neural Network Basics
- Scaling Deep Neural Network Training
- Support Vector Machines

### **Today: Part 2, Unsupervised Learning**

- Non-Negative Matrix Factorization
- Spectral and Markov Clustering
- Sparse Inverse Covariance Matrix Estimation

## Sparse Inverse Covariance Matrix Estimation

---

- **Precision matrix = Inverse covariance matrix**
- **Goal:** Estimating graphical model structure
- “The zeros of a precision matrix correspond to zero partial correlation, a necessary and sufficient condition for conditional independence (Lauritzen, 1996)”
- Sparsity often enforced by regularization
- One algorithm (HP-CONCORD)’s objective function:

$$\underset{\Omega \in \mathbf{R}^{p \times p}}{\text{minimize}} \quad -\log \det(\Omega_D^2) + \text{tr}(\Omega S \Omega) + \lambda_1 \|\Omega_X\|_1 + \frac{\lambda_2}{2} \|\Omega\|_F^2,$$

- $\Omega$  is the sparse inverse covariance matrix we are trying to estimate

# Why care? Finding Direct Associations

---

**Partial Correlation (a.k.a. sparse inverse covariance estimation):  
direct association without confounders**

- Gene Regulatory Network (GRN) estimation
- Joint modeling of SNPs and GRN
- Linkage Disequilibrium (LD) estimation
- Canonical Correlation Analysis (CCA)
- Genome-wide association studies (GWAS)

**Data-driven hypothesis generation!**



**Fig. 1.** Conditionally on the height of snow, the number of snowmen is independent of the intensity of traffic jams. This is represented by a two edges graph.

- Computationally challenging;
- **HP-CONCORD** on distributed memory increases scalability

# HP-CONCORD Algorithm

---

**Algorithm 2** The Cov variant of HP-CONCORD, for computing a sparse estimate of the inverse covariance matrix.

---

**Input:** data matrix  $X \in \mathbf{R}^{n \times p}$ ; tuning parameters  $\lambda_1, \lambda_2 \geq 0$ ; optimization tolerance  $\epsilon > 0$

**Output:** estimate  $\hat{\Omega} \in \mathbf{R}^{p \times p}$  of the underlying inverse covariance matrix  $\Omega^0$

```

1:  $\Omega^{(0)} \leftarrow I$ 
2: Compute  $S \leftarrow \frac{1}{n} X^T X$  ▷ Compute (once) via a distributed dense-dense matrix multiplication
3: Compute  $W^{(0)} \leftarrow \Omega^{(0)} S$  ▷ Compute via a distributed sparse-dense matrix multiplication
4: for  $k = 0, 1, 2, \dots$ 
5:   Form  $(W^{(k)})^T$  ▷ Form via a distributed matrix transpose
6:    $G^{(k)} \leftarrow -(\Omega_D^{(k)})^{-1} + \frac{1}{2}((W^{(k)})^T + W^{(k)}) + \lambda_2 \Omega^{(k)}$  ▷ Use  $W^{(k)}, (W^{(k)})^T$ 
7:    $g(\Omega^{(k)}) \leftarrow -2 \sum_i \log(\Omega_{ii}^{(k)}) + \text{tr}(W^{(k)} \Omega^{(k)}) + \frac{\lambda_2}{2} \|\Omega^{(k)}\|_F^2$  ▷ Use  $(W^{(k)})^T$ ; see text for details
8:   for  $\tau = 1, \frac{1}{2}, \frac{1}{4}, \dots$ 
9:      $\Omega^{(k+1)} \leftarrow \mathcal{S}_{\tau \lambda_1}(\Omega^{(k)} - \tau G^{(k)})$  ▷ Apply the soft-thresholding operator,  $\mathcal{S}_{\tau \lambda_1}$ , in a distributed manner
10:    Compute  $W^{(k+1)} \leftarrow \Omega^{(k+1)} S$  ▷ Compute via a distributed sparse-dense matrix multiplication
11:     $g(\Omega^{(k+1)}) \leftarrow -2 \sum_i \log(\Omega_{ii}^{(k+1)}) + \text{tr}(W^{(k+1)} \Omega^{(k+1)}) + \frac{\lambda_2}{2} \|\Omega^{(k+1)}\|_F^2$  ▷ See text for details
12:    until  $g(\Omega^{(k+1)}) \leq g(\Omega^{(k)}) - \text{tr}((\Omega^{(k)} - \Omega^{(k+1)})^T G^{(k)}) + \frac{1}{2\tau} \|\Omega^{(k)} - \Omega^{(k+1)}\|_F^2$  ▷ See text for details
13: until a stopping criterion is satisfied, using  $\epsilon$ 
14: return the estimate  $\hat{\Omega} \leftarrow \Omega^{(k)}$ 

```

---






- Repeated use of sparse times dense matrix multiplication (SpDM<sup>3</sup>)
- SpDM<sup>3</sup> is the bottleneck by a large margin.

Koanantakool et al. Communication-avoiding optimization methods for distributed massive-scale sparse inverse covariance estimation. In AISTATS, 2018.

# Sparse Matrix times Dense Matrix

Sometimes it pays off to communicate A instead.

How much of these ranges apply to real life NMF scenarios is open question (class project?)



Koanantakool, Penporn, et al. "Communication-avoiding parallel sparse-dense matrix-matrix multiplication." IPDPS, 2016

# HP-CONCORD Advantages

- HP-CONCORD makes fewer assumptions about the data (in particular, no Gaussianity is assumed) compared to competitors
- Thanks to **communication-avoiding matrix multiplication algorithms**, it reaches unprecedented scales

- BigQUIC: previous state-of-the-art
- Obs-K are the other variant of HP-CONCORD algorithm (K: number of nodes)
- Experiment is trying to recover a random graph structure.

